

Research Article

Cultivar-dependent transcriptional responses of lychee to *Peronophythora litchii*: a registered genome-wide analysis

Submission ID f04bfc30-784a-4fe5-9aec-95ed965d4871

Submission Version Initial Submission

PDF Generation 05 Sep 2026 15:19:44 EST by Atypon ReX

Authors

Mr. Eric Zhuang
Corresponding Author
Submitting Author

 [ORCID](https://orcid.org/0009-0001-9050-0214)
<https://orcid.org/0009-0001-9050-0214>

Affiliations
• NYU Langone Health, New York, NY, USA

Files for peer review

All files submitted by the author for peer review are listed below. Files that could not be converted to PDF are indicated; reviewers are able to access them online.

Name	Type of File	Size	Page
Plant_Direct_manuscript.docx	Main Document - MS Word	69.1 KB	Page 3
Figure_1_study_design_and_QC.pdf	Figure	30.4 KB	Page 40
Figure_2_discovery_and_legacy_audit.pdf	Figure	152.3 KB	Page 41
Figure_3_robustness_and_external_evaluation.pdf	Figure	32.0 KB	Page 42
Figure_4_pathway_and_signature_evaluation.pdf	Figure	26.5 KB	Page 43
Figure_5_transcript_usage.pdf	Figure	25.9 KB	Page 44
Figure_6_orthogonal_evidence_and_tiers.pdf	Figure	26.0 KB	Page 45
lychee_unified_manuscript_supplement.zip	Supplementary Material for Review	17.6 MB	Not converted to PDF
Plant_Direct_supporting_figures.pdf	Supplementary Material for Review	1.4 MB	Page 46

**Cultivar-dependent transcriptional responses of lychee to
Peronophythora litchii: a registered genome-wide analysis**

Eric Zhuang

NYU Langone Health, New York, NY, USA

Corresponding author: Eric Zhuang

Email: eric.zhuang@nyulangone.org

ORCID iD: 0009-0001-9050-0214

Article type: Original Research Article

10 Abstract

11 Litchi downy blight, caused by the oomycete *Peronophythora litchii*, is among the most
 12 damaging diseases of lychee (*Litchi chinensis* Sonn.), and public RNA-seq cohorts spanning
 13 cultivars, tissues, and infection time points offer a natural resource for identifying cultivar-
 14 dependent infection responses. Such data are easily over-interpreted, however, when the same
 15 cohort is used both to select candidates and to appear to confirm them. Here we unify two stages
 16 of a transcriptomic reanalysis of the lychee-*P. litchii* interaction into a single account. In the
 17 first, exploratory stage, we reanalyzed leaf RNA-seq data from the cultivars Guiwei and
 18 Yurong1 (GSE201243/PRJNA830488), identified 17 and 117 infection-responsive genes in the
 19 two cultivars, and prioritized 18 candidates with defense-associated annotations and apparently
 20 enriched stress- and hormone-associated promoter elements. In the second, confirmatory stage,
 21 we prospectively registered a discovery and external-evaluation protocol that fixed dataset roles,
 22 statistical models, thresholds, and evidence vocabulary before any external outcome was
 23 examined. A genome-wide negative-binomial cultivar-by-infection interaction analysis of 19,445
 24 expressed genes produced 262 statistical candidates, of which 206 were retained after uniform
 25 mapping and gene-model quality control, 19 remained significant under an independent genome-
 26 wide edgeR quasi-likelihood analysis, and 16 passed the complete registered robustness
 27 procedure. In the prespecified cross-context evaluation against an independent pericarp time
 28 course contrasting a susceptible and a resistant cultivar (PRJNA450886), two genes were
 29 supported and five were directionally contradictory; the internally robust and externally
 30 supported sets did not overlap, an outcome compatible with chance given the set sizes. A formal
 31 audit showed that none of the original 18 exploratory candidates met the newly applied genome-
 32 wide interaction criterion and that two failed uniform-mappability control. Deterministic
 33 evidence integration assigned no gene to the highest confidence tier, placed 12 internally robust

genes, whose annotations include carbohydrate and cell-wall remodeling, lipid-transfer, and chaperone-regulatory functions, in a middle tier, and retired 65 entities, with all null and contradictory outcomes reported. The unified analysis shows that candidates selected for effect-size prominence in a single cohort largely fail a formally specified genome-wide interaction test, delimits exactly which computational claims about this pathosystem the current public data can support, and releases a candidate resource whose evidence boundaries are explicit. Experimental perturbation remains necessary for any causal conclusion.

Keywords: *Litchi chinensis*; *Peronophythora litchii*; RNA-seq; cultivar-by-infection interaction; preregistered analysis; cross-context evaluation

1. Introduction

Lychee (*Litchi chinensis* Sonn.) is a subtropical fruit crop of substantial economic importance in southern China, Southeast Asia, and a growing set of production regions worldwide [1]. Among its diseases, litchi downy blight, caused by the oomycete *Peronophythora litchii* (historically also treated as *Phytophthora litchii* [2]), is particularly damaging: the pathogen attacks young leaves, inflorescences, and fruit, produces water-soaked brown lesions, and drives severe postharvest decay [2,3]. Because chemical control is costly and resistance breeding depends on understanding how host genotypes differ in their response to infection, the molecular basis of cultivar-level variation in the response to *P. litchii* is a question of both agronomic and biological interest.

The resources needed to address this question computationally have accumulated rapidly. High-quality genome assemblies are available for lychee [4], and several groups have profiled the transcriptional response to *P. litchii* infection. Sun et al. compared a susceptible and a resistant cultivar across a pericarp infection time course and reported that the two genotypes deploy distinct early response programs [5]. Liu et al. described calcium-dependent protein kinase

(LcCDPK) family members induced by *P. litchii* inoculation [6], and Li et al. showed that the pathogen RXLR effector PIAvh202 promotes virulence by destabilizing a host ethylene-biosynthesis enzyme [7]. A pathogen-side transcriptome catalogued the *P. litchii* pathogenicity arsenal [2]. In parallel, public repositories now hold RNA-seq series covering resistant and susceptible cultivars, leaf and fruit tissues, messenger and small RNA, and multiple time points. These series offer an opportunity, in principle, to distinguish genome-wide cultivar-dependent infection responses from responses common to all genotypes. Realizing that opportunity is statistically demanding. The quantity of interest is not differential expression within a cultivar but the cultivar-by-infection interaction: the difference between infection responses across genotypes. Interaction effects are estimated with less precision than main effects, the deposited designs are small (typically three libraries per cultivar-treatment cell), and a genome-wide search requires an explicit multiplicity family. Analyses that instead select a handful of strongly responding genes and then characterize them in the same cohort conflate selection with confirmation: large observed fold changes in a small experiment are a mixture of true signal and noise, and features chosen for extremity will regress toward null in any independent assessment. These concerns apply directly to the present pathosystem, in which most published candidate genes derive from single-cohort analyses without external evaluation. The first stage of the present work followed the same pattern. In an exploratory reanalysis of the GSE201243 leaf series [8] (Guiwei and Yurong1, mock versus inoculated, 24 h post-inoculation), we identified 17 significant infection-responsive genes in Guiwei and 117 in Yurong1, prioritized 18 candidates with defense-plausible annotations (among them jacalin-like lectin, pathogenesis-related and thaumatin/osmotin-like proteins, a heat-shock protein, lipid-transfer proteins, and calcium-binding proteins), and found that a set of stress- and hormone-associated *cis*-regulatory elements was enriched in their promoters relative to a randomized

82 background. That analysis was framed as hypothesis-generating, and critical review of it
 83 converged on three substantive criticisms: the cultivar-by-treatment comparison was an effect-
 84 size contrast over pre-selected genes rather than a genome-wide interaction test with false-
 85 discovery-rate control; no independent dataset was used for validation; and promoter-element
 86 counting, even against a randomized background, is weak evidence of regulatory function. We
 87 regard all three criticisms as correct, and they motivated the confirmatory stage reported here.
 88 This paper therefore presents the complete arc of the project in one place. We prospectively
 89 registered a discovery and external-evaluation protocol that fixed dataset roles, models,
 90 thresholds, empty-result behavior, and the vocabulary used to describe evidence before any
 91 external outcome was opened. Because every cohort is public, this staged unlock is a
 92 computational safeguard enforced by the workflow rather than blinding in the experimental
 93 sense, and we document the registration chronology in the Methods. Within the prespecified
 94 discovery cohort we fitted a genome-wide negative-binomial interaction model, applied uniform
 95 mapping and gene-model quality control, and subjected the resulting candidates to a battery of
 96 internal robustness checks spanning independent quantification, an alternative statistical
 97 framework, filter sensitivity, and leave-one-library-out refits. We then evaluated the frozen
 98 candidates (throughout, *frozen* denotes a computed result fixed at a recorded checkpoint and not
 99 modified afterward), a frozen pathway set, and a frozen signed expression signature in
 100 independent public cohorts whose roles (primary cross-context evaluation, generic transfer,
 101 orthogonal modality, or exploratory) were fixed in advance. Orthogonal evidence classes
 102 (sequence annotation, promoter motifs, small-RNA coherence, and accession-aware literature)
 103 were kept separate from expression support, and a deterministic tier system integrated all layers
 104 while explicitly permitting an empty top tier. Finally, we audited the 18 exploratory candidates
 105 from the first stage against the genome-wide analysis. The result is a candidate resource for

lychee downy blight whose every claim carries its evidence class, together with a case study in how much, and how little, of a plausible exploratory analysis survives formal statistical scrutiny.

2. Results

2.1 A two-stage design with prespecified dataset roles

The study comprises an exploratory stage completed first and a confirmatory stage whose protocol was registered before its analyses began (Figure 1A; Table 1). Five public cohorts enter the confirmatory stage in fixed, non-interchangeable roles. PRJNA830488 (GSE201243), the leaf series for Guiwei and Yurong1 at 24 h, is the sole discovery cohort; it is also the dataset analyzed in the exploratory stage. PRJNA450886, a pericarp time course contrasting the susceptible cultivar Guiwei with the resistant cultivar Heiye at 6, 24, and 48 h [5], serves as the primary external evaluation. Because it differs from the discovery cohort in tissue, resistant comparator, and time structure, we describe positive outcomes in it as *cross-context support* rather than replication. PRJNA922966 (Feizixiao leaf and fruit, mock versus infected) tests generic infection-response transfer but cannot address the cultivar interaction; PRJNA922965 provides small-RNA libraries as an orthogonal modality; and PRJNA1090613 (Guiwei versus SFZ leaf) is exploratory only, because its time and resistance metadata could not be resolved. Under the registered firewall, no external outcome could alter the frozen discovery membership, signature weights, or thresholds, and the protocol explicitly allowed the top evidence tier to be empty.

[Insert Figure 1 near here]

Table 1. Locked dataset roles and eligibility.

Accession	Design	Locked role	Evidence class
PRJNA830488 (GSE201243)	Guiwei/Yurong1 leaf; mock/infected; 24 h; 3 libraries per cell	Discovery	Conditional inferential

Accession	Design	Locked role	Evidence class
PRJNA450886	Guiwei/Heiye pericarp; mock/infected; 6, 24, 48 h; 3 libraries per cell	Primary external evaluation	Cross-context
PRJNA922966 (GSE222651)	Feizixiao leaf/fruit; mock/infected; 24 h; 3 libraries per cell	Generic infection transfer	Cross-context
PRJNA922965 (GSE222650)	Feizixiao leaf/fruit small RNA; mock/infected; 24 h; 3 libraries per cell	Orthogonal modality	Orthogonal
PRJNA1090613 (GSE262200)	Guiwei/SFZ leaf; mock/infected; 3 libraries per cell	Exploratory only	Exploratory

2.2 Stage 1: exploratory candidate discovery in GSE201243

The exploratory stage analyzed the 12 discovery libraries with per-cultivar differential-expression contrasts (infected versus mock) and identified 17 significant genes in Guiwei (15 up-regulated, 2 down-regulated) and 117 in Yurong1 (39 up-regulated, 78 down-regulated) at Benjamini–Hochberg adjusted $P < 0.05$, suggesting a broader transcriptional response in Yurong1. From these lists, 18 genes were prioritized by combining statistical significance in at least one cultivar, large infection-associated fold changes, and annotations plausibly related to defense: among them a jacalin-like lectin (LITCHI028104), a pathogenesis-related protein (LITCHI011394), osmotin/thaumatin-like proteins (LITCHI009301, LITCHI002793), an HSP70-family chaperone (LITCHI012388), lipid-transfer proteins (LITCHI019183, LITCHI018043), an EF-hand calcium-binding protein (LITCHI019519), and a phytosulfokine precursor (LITCHI007102). A selected-gene contrast of infection responses between cultivars nominated LITCHI019519, LITCHI001510, LITCHI028401, and LITCHI017676 as Guiwei-skewed and LITCHI028104 and LITCHI019183 as Yurong1-skewed candidates.

The exploratory stage also examined 2-kb upstream promoter regions of the 18 candidates. Against 100,000 randomized GC-matched promoter sets, four *cis*-element classes appeared enriched: anaerobic-response elements (ARE; 46 observed versus 24.6 expected, $q = 0.0003$), methyl-jasmonate-responsive motifs (54 versus 38.4, $q = 0.018$), TCA elements (9 versus 0.19, q

= 0.00003), and TC-rich repeats (10 versus 0.15, $q = 0.00003$). ABRE and MYB-binding-site motifs, in contrast, were frequent but not enriched. All of these results were obtained by selecting and characterizing genes within a single cohort; they carry no genome-wide multiplicity control and no external evaluation, and we treat them below strictly as hypotheses submitted to the confirmatory stage.

2.3 Discovery-cohort quality under the registered protocol

The confirmatory pipeline combined a joint host and pathogen reference, fastp trimming, STAR two-pass alignment, featureCounts gene counting, and independent Salmon quantification. Under this pipeline, all 12 discovery libraries passed the fixed technical gates of at least 10 million surviving read pairs and at least 40% uniquely aligned reads (Figure 1B). Read survival after trimming ranged from 99.56% to 99.66%, unique-alignment rates from 84.44% to 86.22%, and Salmon mapping rates from 80.77% to 82.56%. Principal-component analysis of the transformed counts separated libraries by cultivar on the first component (54.2% of variance) and by infection status on the second (16.1%), with no library isolated from its replicate group (Figure 1C). Because source-tree, pooling, and extraction independence are not reported for the deposited libraries, all inferential statements below are conditional on deposited-library independence, a limitation we return to in the Discussion.

2.4 Genome-wide cultivar-by-infection discovery

The primary discovery model fitted all expressed genes with the design $\sim$ cultivar + treatment + cultivar:treatment, with Guiwei and mock as reference levels, so that the interaction coefficient estimates how much the infected-minus-mock response in Yurong1 exceeds that in Guiwei. Of 19,445 expressed genes, 262 met the registered discovery threshold of genome-wide Benjamini–Hochberg $q < 0.05$ with an absolute interaction log2 fold change of at least log2(1.5) (Figure

2A); 277 genes met the significance criterion alone and 5,370 met the effect criterion alone. All interaction estimates reported and plotted here are unshrunk maximum-likelihood values; apeglm shrinkage was used only to define the frozen signature weights. Uniform quality control then removed 31 candidates for exon-level mappability failure and 25 for gene-model ambiguity, freezing 206 genes as statistical candidates (Figure 2C). Because these quality filters were applied after significance testing, we verified that the ordering is not load-bearing: restricting the testing universe to the 13,602 genes that pass uniform quality control and recomputing the Benjamini–Hochberg adjustment over that universe rediscovers all 206 frozen candidates and adds 12 further genes, so the frozen list is a conservative subset of the QC-first analysis. A post hoc composite-null sensitivity analysis asked the stricter question whether $|\beta|$ exceeded $\log_2(1.5)$, rather than testing $\beta = 0$ and filtering the estimate afterward. Eight of the 262 statistical candidates passed the DESeq2 composite-null test at $q < 0.05$; the corresponding counts were 7 of 206 QC-retained candidates, 7 of 19 edgeR-gated candidates, and 7 of 16 fully robust candidates. An apeglm false-sign-or-small analysis was less severe, with $s < 0.05$ for 149, 113, 18, and 15 genes in those four nested sets, respectively. These are post hoc sensitivity results, not a redefinition of the registered discovery set. Positive interaction effects (stronger infection response in Yurong1) and negative effects (stronger response in Guiwei) were both well represented. At the pathway level, six Plant Reactome pathways transported through one-to-one *Oryza* orthology met the frozen discovery threshold: activation of the pre-replication complex, assembly of the pre-replication complex, DNA replication initiation, maturation, and circadian rhythm, all with positive enrichment, and jasmonic acid signaling with negative enrichment. Together with the genes, these six pathways, the apeglm-shrunk signed signature weights over the 206 genes, and the conditional transcript-usage results (Section 2.9) were frozen before any external data were opened.

[Insert Figure 2 near here]

2.5 None of the 18 exploratory candidates met the genome-wide interaction criterion

Because the exploratory candidates were selected before any interaction P -values existed, the registered protocol included a formal audit of all 18 against the genome-wide analysis (Figure 2B; Table 2). None met the genome-wide interaction threshold. We then performed the estimand-aligned follow-up requested in review: separate confirmatory DESeq2 infected-versus-mock fits within each cultivar, with Benjamini–Hochberg adjustment over the same 19,445-gene universe in each contrast (Supplementary S14). Thirteen of the 18 legacy genes were significant in Guiwei and 16 of 18 in Yurong1 at $q < 0.05$; all 13 Guiwei genes and 15 of the 16 Yurong1 genes also exceeded $|\log_2\text{FC}| \geq \log_2(1.5)$. Thus most of the original within-cultivar response claims survive the confirmatory pipeline even though none supports a genome-wide cultivar-by-infection interaction. The strongest interaction-audit candidate, LITCHI001510, reached $q = 0.080$; the EF-hand calcium-binding gene LITCHI019519, whose selected-gene contrast of -3.4 had appeared most dramatic in the exploratory stage, reached only $q = 0.106$ despite a re-estimated interaction effect of -3.33 ; and the median genome-wide interaction q across the 18 candidates was 0.47 . Two candidates, the osmotin/thaumatin-like gene LITCHI009301 and LITCHI028401, additionally failed uniform exon-mappability control, meaning their interaction estimates cannot be attributed unambiguously to a single locus. The audit therefore does not show that these genes are biologically inert. It shows that within-cultivar responsiveness and cultivar-dependent response are distinct claims, and only the latter fails here.

Table 2. Formal genome-wide audit of the 18 exploratory-stage candidates. Interaction effects are re-estimated genome-wide values (Yurong1 response minus Guiwei response, $\log_2$ scale). No candidate met the registered discovery threshold (genome-wide $q < 0.05$ and $|\log_2\text{FC}| \geq \log_2(1.5)$).

Gene	Exploratory annotation	Interaction $\log_2\text{FC}$ (SE)	Genome-wide q	Uniform mappability
LITCHI001510	Poorly characterized	-2.83 (0.91)	0.080	Pass

Gene	Exploratory annotation	Interaction log2FC (SE)	Genome-wide q	Uniform mappability
	responsive gene			
LITCHI019519	EF-hand calcium-binding protein	-3.33 (1.14)	0.106	Pass
LITCHI019299	S-adenosylmethionine-related protein	-1.21 (0.47)	0.175	Pass
LITCHI007102	Phytosulfokine precursor	-0.74 (0.32)	0.238	Pass
LITCHI019183	Plant lipid-transfer protein	+0.99 (0.45)	0.273	Pass
LITCHI028401	Selected interaction candidate	-2.41 (1.13)	0.281	Fail
LITCHI012388	Heat-shock protein 70	-0.90 (0.49)	0.381	Pass
LITCHI001552	Cyclase-like stress-related protein	+0.92 (0.55)	0.430	Pass
LITCHI001512	Poorly characterized responsive gene	-1.01 (0.63)	0.457	Pass
LITCHI017676	S-adenosylmethionine-related protein	-2.36 (1.53)	0.478	Pass
LITCHI002550	Cytochrome P450	+0.76 (0.96)	0.762	Pass
LITCHI028104	Jacalin-like lectin	+1.39 (1.89)	0.785	Pass
LITCHI002793	Osmotin/thaumatin-like protein	+0.41 (0.55)	0.786	Pass
LITCHI017995	Metalloendoproteinase	+0.48 (0.73)	0.812	Pass
LITCHI009301	Osmotin/thaumatin-like protein	-0.70 (1.08)	0.816	Fail
LITCHI025738	Cyclase-like stress-related protein	-0.22 (0.62)	0.908	Pass
LITCHI011394	Pathogenesis-related protein	-0.25 (0.98)	0.940	Pass
LITCHI018043	Plant lipid-transfer protein	-0.10 (0.52)	0.954	Pass

217

218 2.6 Sixteen genes and one pathway passed all internal robustness gates

219 Each of the 206 frozen genes was required to retain its interaction direction and significance
220 under four predeclared perturbations of the analysis: independent Salmon-based quantification
221 (196 of 206 passed), edgeR quasi-likelihood testing in place of the primary framework (19
222 passed), a counts-per-million expression filter (205 passed), and all twelve leave-one-library-out
223 refits (125 passed). Eighteen genes passed all four gates jointly, and 16 remained after an
224 additional observed mapping-sensitivity gate that excluded candidates whose effect depended on
225 ambiguously mapped reads (Figure 3A). The alternative-framework gate was by far the most

stringent, and the frozen genome-wide edgeR results explain why. The two frameworks agree almost perfectly on estimation: interaction effects correlate at Pearson $r = 0.995$ across all 19,445 genes and $r = 0.999$ across the 206 candidates, sign agreement among the candidates is 100%, and every candidate has nominal edgeR $P < 0.05$. The disagreement is confined to significance calibration. The registered gate required genome-wide edgeR $q < 0.10$ with direction agreement, and edgeR's quasi-likelihood test with robust dispersion estimation admits only 24 genes genome-wide at that threshold (10 at $q < 0.05$), against 262 for the primary Wald analysis at $q < 0.05$. With three libraries per cell, the two dispersion-estimation strategies imply very different genome-wide significance sets even while agreeing on every point estimate. We therefore present the results as a hierarchy: 262 statistical candidates, 206 retained after quality control, 19 significant under both frameworks, and 16 passing the complete registered procedure; genes surviving both frameworks constitute the discovery set's stable core, and the wider 206-gene list should be read as one framework's candidates rather than as settled discoveries. Of the six frozen pathways, only circadian rhythm passed its four conjunctive gates (competitive and rotation-based testing, leading-edge deletion, and an expression- and length-matched null); the other five all failed the rotation-based self-contained test, and the maturation set additionally failed leading-edge deletion, indicating enrichment that was not stable across testing frameworks.

[Insert Figure 3 near here]

2.7 The prespecified external evaluation supported two genes and contradicted five

The primary external estimand, fixed before unlock, was the resistant-minus-susceptible cultivar-by-infection interaction at 24 h in the PRJNA450886 pericarp time course, tested for each frozen candidate with Benjamini–Hochberg adjustment across the full candidate-by-contrast family, an effect threshold of $\log_2(1.5)$, confidence intervals excluding zero, and required agreement with the discovery direction. Of the 206 frozen genes, 184 were measurable in the external cohort.

Two met every criterion and are cross-context supported: LITCHI005518, annotated as an alpha-xylosidase involved in xyloglucan remodeling of the cell wall, and LITCHI001963, an uncharacterized gene with a one-to-one *Oryza* ortholog. Five genes, including a beta-glucosidase and a Fe(II)/2-oxoglutarate-dependent dioxygenase, passed the magnitude and significance criteria with the *opposite* direction and are recorded as contradictory; 177 were unsupported at the prespecified thresholds and 22 not testable (Figure 3C; Table 3). At the level of continuous estimates, discovery and external interaction effects across the 184 measurable candidates were uncorrelated (Pearson $r = -0.05$, 95% CI -0.19 to 0.10 ; Spearman $\rho = -0.04$), with sign concordance of 51.6% (95 of 184), indistinguishable from chance.

The two supported genes are not among the 16 internally robust genes, and none of the 16 internally robust genes achieved external support (Figure 3B). This non-overlap is, by itself, nearly uninformative: for sets of 16 and 2 genes drawn from 206 candidates, the expected overlap under independence is 0.16 genes and the probability of zero overlap is 0.85, so the disjoint sets neither demonstrate nor refute an association between internal stability and cross-context transport. The absence of continuous concordance reported above is the more informative observation. We emphasize that the external cohort differs in tissue (pericarp versus leaf), resistant comparator (Heiye versus Yurong1), and design (time course versus single time point), so absence of support is expected wherever a response is tissue- or genotype-specific, and support that is observed indicates transport across contexts rather than replication of the original experiment. Because Yurong1 and Heiye have not been placed on a common resistance scale relative to Guiwei, we interpret the primary external test as cross-cultivar response transport rather than validation of a resistance phenotype.

Table 3. Cross-context supported and contradictory genes in the primary external evaluation (PRJNA450886, pericarp, 24 h). Discovery effects are Yurong1-minus-Guiwei interaction log2 fold changes in leaf; external effects are Heiye-minus-Guiwei interaction log2

fold changes in pericarp with 95% confidence intervals. All q -values are Benjamini–Hochberg adjusted (discovery: genome-wide; external: across the frozen candidate-by-contrast family).

Gene	Annotation	Discovery log2FC	Discovery q	External log2FC (95% CI)	External q	Outcome
LITCHI005518	Alpha-xylosidase 1	+0.64	0.025	+0.67 (+0.32 to +1.02)	0.007	Supported
LITCHI001963	Uncharacterized; one-to-one <i>Oryza</i> ortholog	−1.41	0.024	−0.94 (−1.43 to −0.46)	0.007	Supported
LITCHI030761	Caffeoylshikimate esterase	−0.92	0.040	+1.31 (+0.78 to +1.85)	0.0002	Contradictory
LITCHI029534	Thiamine thiazole synthase 2, chloroplastic	+1.48	0.002	−1.04 (−1.61 to −0.48)	0.010	Contradictory
LITCHI021860	Phosphatidylinositol 4-kinase gamma 2	−0.65	0.032	+0.97 (+0.39 to +1.54)	0.029	Contradictory
LITCHI008313	Beta-glucosidase 44	+0.90	0.045	−1.12 (−1.81 to −0.44)	0.034	Contradictory
LITCHI001700	Fe(II)/2-oxoglutarate-dependent dioxygenase 4	+3.70	0.033	−2.00 (−3.24 to −0.77)	0.035	Contradictory

2.8 Pathway and signature transport was weak and strongly context-dependent

None of the six frozen pathways, including the internally robust circadian-rhythm pathway, passed the external pathway gates in PRJNA450886 (Figure 4A). The frozen signed signature over the 206 discovery genes behaved in a strongly context-dependent way (Figure 4B). At the primary 24-h contrast the signature score was indistinguishable from zero (estimate 0.109, 95% CI −0.079 to 0.297, $q = 0.349$), and at 48 h it was significantly *negative* (−0.240, 95% CI −0.428 to −0.052, $q = 0.043$), a sign reversal relative to discovery. In the generic-transfer study the signature tracked infection positively in leaf (0.107, $q = 0.0015$) but negatively in fruit (−0.181, $q < 0.001$), with a significant leaf-minus-fruit interaction (0.289, $q < 0.001$). The signature therefore carries reproducible infection-associated information, but its sign and magnitude depend on tissue and time in ways that preclude its use as a context-free resistance marker. The exploratory cohort's estimate is reported separately in Figure S3 for completeness and is not confirmatory.

[Insert Figure 4 near here]

2.9 Conditional transcript-usage discovery remained internal

Differential transcript usage was analyzed under a conditional gate registered in advance: stage-wise testing had to yield interpretable gene- and transcript-level error control, and no post hoc substitute was permitted if the gate failed. The gate passed, and 225 transcript-usage events across 152 genes were frozen from 15,790 tested transcripts (Figure 5A). External follow-up preserved the role separation of the cohorts (Figure 5B): in the primary cross-context study, 125 events were measurable but unsupported at the prespecified thresholds and 100 were not testable; in the generic-transfer study, annotation incompatibilities left all 225 untestable; and in the exploratory cohort, 4 events were supported and 5 contradictory; we report these outcomes descriptively and do not count them as primary external support, because that cohort's metadata could not be standardized. Transcript-level rewiring between cultivars therefore remains an internally supported, externally unconfirmed layer of the resource.

[Insert Figure 5 near here]

2.10 Orthogonal evidence: informative nulls and reference-ecosystem gaps

The registered orthogonal evidence classes returned uniformly negative or non-testable results, and we report them as such rather than converting them into support. Computational annotation assigned family-level functional labels to 150 of the 206 candidates and left 56 unannotated, but no candidate met the registered high-confidence rule requiring two independent evidence classes with at least 70% supported coverage and no architecture conflict, largely because precomputed InterPro domain architectures were unavailable for all 206 proteins in this non-model species (Figure 6A). Registered promoter-motif discovery tested 927 JASPAR plant position-weight matrices in strand-aware 1-kb and 2-kb windows against 100 expression- and GC-matched genomic background sets. No motif passed in more than 4 of 100 backgrounds against a required

80, and none passed FIMO sensitivity (Figure 6B). Small-RNA coherence was not testable because the frozen PmiREN reference contained no exact *Litchi* entries.

A separate post hoc controlled comparison held the published 18-promoter totals, six explicitly reported exact-element strings plus one standard PlantCARE TCA-element variant, scoring rule, and multiplicity correction fixed while changing only the background (Supplementary S15). The 100,000-set randomized-GC rerun closely reproduced the exploratory result. Against 100 expression- and GC-matched genomic promoter sets, ARE lost significance, ABRE gained it, and MeJA-responsive, TCA, and TC-rich classes remained significant: four of seven classes passed under either strategy, but the identity of one class changed. Background construction therefore affected the result but did not by itself explain the exploratory/registered discrepancy. This comparison is provenance-limited: the archived exploratory package omitted the observed-motif input TSV, and exact recounting from the current canonical 18 promoters reproduced the published total only for ARE. Accordingly, it cannot rescue a functional regulatory claim or establish which foreground representation generated the other published totals. Published studies reusing any analyzed accession, including the primary external study itself [5], remain registered as prior interpretation rather than independent support.

[Insert Figure 6 near here]

2.11 Deterministic integration: an empty top tier, twelve Tier B genes, and visible retirements

The registered tier rules integrate all evidence layers without discretion. Tier A required discovery, complete internal robustness, primary cross-context support, clean mapping and annotation, and at least one attributable orthogonal class; Tier B required discovery and full internal robustness with partial or non-testable external evidence; Tier C was reserved for a robust, externally supported pathway. Applied to the 268 frozen entities (262 statistical gene

candidates and 6 pathways), the rules assigned no entity to Tier A or Tier C, 12 genes to Tier B,
 191 entities to exploratory status, and 65 to retirement (Figure 6C). The retirements comprise the
 56 statistical candidates removed by uniform mapping or gene-model control, the 5 genes
 contradicted for cross-context directional transport, and 4 candidates with candidate-level
 observed-mapping or annotation failures. These categories differ in kind, and the released tables
 preserve the distinction: mapping and annotation failures undermine the discovery itself, whereas
 an opposite direction in a different tissue and cultivar contrast may reflect genuine context
 specificity rather than an invalid discovery in the original context. Four internally robust genes,
 including a probable RPP8-like disease-resistance protein, were measurable but unsupported at
 the prespecified thresholds in the primary external evaluation and were therefore held at
 exploratory status rather than promoted. The two cross-context-supported genes likewise remain
 exploratory because they failed internal robustness. All contradictory and failed outcomes are
 released alongside the positive ones. We also note that Tier A was difficult to reach by
 construction in the current reference ecosystem: with the small-RNA class structurally
 untestable, no motif able to pass the family-wise rule, accession-linked literature excluded as
 independent support, and the two-class annotation rule unattainable without precomputed
 domain architectures, the orthogonal-class requirement could not realistically be satisfied by any
 gene. The empty top tier therefore reflects this structural constraint as much as the expression
 evidence itself.

The 12 Tier B genes are listed in Table 4. Nine were measurable in the primary external cohort,
 and all nine were direction-consistent with discovery, although none met the full support criteria.
 Figure S1 exposes the normalized count of every individual discovery library for these 12 genes,
 the 2 cross-context-supported genes, and legacy highlights LITCHI001510 and LITCHI019519;
 every cultivar-treatment cell contains all three deposited replicates. Their annotations cluster

around carbohydrate and cell-wall metabolism (a phosphomannomutase/phosphoglucomutase, a chloroplastic 1,4- α -glucan-branching enzyme, an ERD6-like sugar transporter, a serine carboxypeptidase), membrane transport (a WAT1-related protein), lipid signaling (a DIR1-like lipid-transfer protein, the *Arabidopsis* ortholog of which is required for systemic acquired resistance), programmed-cell-death regulation (a BAG-family chaperone regulator), specialized metabolism (cytochrome P450 81Q32), RNA metabolism (ribonuclease J, exonuclease 1), chromatin-associated regulation (a homeobox-DDT protein), and one entirely unannotated gene.

Table 4. The twelve Tier B genes: internally robust genome-wide statistical candidates with direction-consistent but sub-threshold, or non-measurable, primary external evidence. Discovery effects are interaction log₂ fold changes (Yurong1 response minus Guiwei response); positive values indicate a stronger infection response in Yurong1. External columns give the primary 24-h pericarp estimate with 95% confidence interval and family-adjusted *q*; NM denotes genes not measurable in the external cohort.

Gene	Family-level annotation	Discovery log ₂ FC	Genome-wide <i>q</i>	External log ₂ FC (95% CI)	External <i>q</i>
LITCHI014613	Sugar transporter ERD6-like 16	+3.69	8.7×10^{-6}	+0.60 (−0.04 to +1.25)	0.27
LITCHI015982	WAT1-related protein	+3.43	7.8×10^{-7}	+0.77 (+0.07 to +1.48)	0.18
LITCHI005805	Phosphomannomutase/phosphoglucomutase	+2.07	5.5×10^{-8}	NM	—
LITCHI028926	Exonuclease 1	+1.90	5.7×10^{-5}	+0.32 (−0.63 to +1.28)	0.83
LITCHI016077	Serine carboxypeptidase 24	+1.81	4.4×10^{-5}	+0.23 (−1.05 to +1.51)	0.89
LITCHI013915	1,4- α -glucan-branching enzyme 3, chloroplastic	+1.58	4.4×10^{-5}	+0.53 (+0.10 to +0.95)	0.13
LITCHI021546	Homeobox-DDT domain protein RLT3	+1.31	6.6×10^{-5}	+0.12 (−0.30 to +0.54)	0.83
LITCHI022605	Ribonuclease J	+0.74	2.9×10^{-4}	+0.14 (−0.34 to +0.62)	0.83
LITCHI028717	BAG-family molecular chaperone regulator 4	−0.91	1.8×10^{-5}	−1.20 (−2.52 to +0.12)	0.29
LITCHI008721	Unannotated protein	−1.91	7.9×10^{-11}	NM	—
LITCHI015740	Putative lipid-transfer protein DIR1	−1.96	1.2×10^{-7}	−0.21 (−1.10 to +0.69)	0.88
LITCHI010877	Cytochrome P450 81Q32	−2.77	2.1×10^{-7}	NM	—

2.12 Parametric simulations quantify the interaction effects detectable with three libraries per cell

The post hoc power analysis simulated negative-binomial counts from the fitted gene-wise means and dispersions while preserving the four-cell discovery design and the 24-h four-cell external design (Figure S2; Supplementary S18). At each of 12 absolute interaction effects from 0 to 4 log₂ units, 100 simulations were run sequentially. Genome-wide discovery used 19,445 tests with 262 injected targets per simulation, matching the observed statistical-candidate prevalence; the external analysis adjusted within the 177 frozen genes that passed the simulation's count and dispersion validity filters, with 2 injected targets matching the number of directionally supported genes. Seven additional genes were measurable in the observed 184-gene external contrast but were not fit-eligible for simulation, so the simulated multiplicity family is reported explicitly rather than equated with the observed family. The overall interpolated effect required for 80% detection was 2.44 log₂FC (5.4-fold) in discovery and 2.11 log₂FC (4.3-fold) in the external candidate-family analysis. At a true |interaction log₂FC| of 1.5, overall detection probabilities were 62.7% and 74.0%, respectively. These values are conditional on fitted dispersions treated as known, deposited-library independence, and the simulated non-null prevalence; they quantify the low-power regime but cannot distinguish context dependence from absence of effect for any particular unsupported gene.

3. Discussion

The central finding of this unified analysis is methodological and biological at once: in the lychee–*P. litchii* pathosystem, the genes that look most impressive in a single-cohort exploratory analysis and the genes that survive formal genome-wide inference are almost entirely different sets. The 18 exploratory candidates were chosen for large fold changes and defense-plausible

400 annotations, yet none reached genome-wide significance for the cultivar-by-infection interaction,
 401 and two could not even be attributed unambiguously to a single locus. Conversely, the 16 genes
 402 that passed every internal robustness gate, and the 12 that reached Tier B, are dominated not by
 403 canonical pathogenesis-related families but by carbohydrate and cell-wall metabolism, transport,
 404 lipid signaling, and RNA metabolism, including one gene with no annotation at all. Neither
 405 observation would have been visible without running both stages, and together they quantify a
 406 selection effect that is likely to operate in any small-replicate transcriptomic study of this kind.
 407 One qualification applies throughout: the audit changed not only the statistical machinery but, for
 408 many of the 18 genes, the scientific question, from within-cultivar responsiveness to cultivar-
 409 dependent response, so the comparison measures survival under a stricter and partly different
 410 criterion rather than direct refutation of the original claims.

411 Why did the exploratory candidates fail the interaction audit? Three mechanisms, all anticipated
 412 by the reviewers of the exploratory report, are sufficient to explain the outcome. First, selection
 413 on observed fold change in a three-replicate contrast preferentially captures genes with high
 414 sampling variance, and their apparent effects regress when the full uncertainty is modeled:
 415 LITCHI019519's selected-gene contrast of -3.4 looked decisive, but its genome-wide interaction
 416 test carried a standard error of 1.14 and a q -value of 0.106. Second, per-cultivar differential
 417 expression followed by informal comparison is not an interaction test; genes can differ
 418 dramatically in within-cultivar significance while their responses do not differ significantly
 419 between cultivars. Third, two candidates failed uniform-mappability control, a failure mode
 420 invisible to any analysis that trusts a single quantification path. None of this implies the
 421 exploratory candidates are biologically irrelevant; several retain large direction-consistent
 422 estimates and remain reasonable targets for quantitative PCR. Their evidentiary status, however,
 423 is now stated correctly: unconfirmed hypotheses, not validated markers.

424 The annotations of the genes that survive scrutiny suggest a recognizable, although still
 425 hypothetical, theme. Annotations related to cell-wall and sugar metabolism appear at every
 426 supported evidence level: among the Tier B genes, a
 427 phosphomannomutase/phosphoglucomutase, a starch-branching enzyme, an ERD6-like sugar
 428 transporter, and a serine carboxypeptidase; and among the two cross-context-supported genes, an
 429 alpha-xylosidase acting on xyloglucan. We emphasize that this is a qualitative clustering of
 430 family-level annotations over small gene sets, assembled from different genes at different
 431 evidence levels, and not a formal over-representation result; no enrichment test was performed,
 432 and several annotations are provisional. The pattern is nonetheless consistent with the emphasis
 433 on sugar metabolism and cell-wall reinforcement in prior comparisons of resistant and
 434 susceptible lychee cultivars [5], and with the general role of apoplastic sugar status and wall
 435 remodeling in oomycete resistance. The Tier B set also contains a DIR1-like lipid-transfer
 436 protein, whose *Arabidopsis* counterpart is a mobile signal in systemic acquired resistance, and a
 437 BAG-family chaperone regulator; BAG proteins act as HSP70 co-chaperones within the
 438 chaperone networks that shape plant immune responses and immune-associated cell death [9].
 439 The strong negative discovery enrichment of jasmonic acid signaling, although it did not survive
 440 external evaluation, is likewise a plausible axis of cultivar difference given the jasmonate-
 441 associated promoter elements suggested in the exploratory stage and the known antagonism
 442 between jasmonate and salicylate defense programs. We advance all of these as prioritized,
 443 testable hypotheses; the tier system exists precisely to prevent us from advancing them as more.
 444 The external evaluation delivers a second structural lesson. The internally robust and externally
 445 supported sets did not overlap, but this disjunction carries almost no information on its own: with
 446 16 and 2 genes selected from 206 candidates, the expected overlap under independence is 0.16
 447 and zero overlap occurs with probability 0.85. The informative observations are continuous ones:

discovery and external interaction effects were uncorrelated across the 184 measurable candidates, the frozen signature reversed sign between 24 h and 48 h in the external time course, and the same signature tracked infection in opposite directions in leaf and fruit of the transfer cohort. The straightforward reading is that the cultivar-dependent component of the early infection response is substantially tissue- and time-specific, so that a candidate's failure to transport from leaf at 24 h to pericarp at 48 h is expected behavior rather than evidence against the discovery. This has a practical corollary for breeding applications: expression-based resistance markers derived from one tissue and time point should not be assumed portable, and marker development for this pathosystem will require validation in the tissue and developmental window where deployment is intended.

The orthogonal layers contribute mainly disciplined null results, and we consider their inclusion one of the more useful aspects of the resource. The post hoc controlled PlantCARE comparison now isolates the background rule as far as the archived inputs permit. It shows that genomic expression/GC matching changes which class passes (ARE is replaced by ABRE) but does not eliminate the MeJA, TCA, or TC-rich signals when the published foreground totals are held fixed. Background choice is therefore not a sufficient explanation for the failure of the registered JASPAR analysis; the changed foreground, motif representation, scoring, and robustness rule remain consequential. Because the original observed-motif TSV was not archived and most published totals cannot be reproduced from the current canonical promoters, the controlled result is a provenance audit rather than evidence of regulatory function. The annotation and small-RNA outcomes likewise expose reference-ecosystem gaps (no precomputed InterPro architectures and no *Litchi* entries in the frozen miRNA reference) that bound what any computational study of this species can currently claim.

471 Several limitations qualify all of the above. The deposited designs are small, with three libraries
472 per cell, and the independence of those libraries at the level of source trees, pooling, and
473 extraction is unreported, so our q -values are conditional on deposited-library independence. The
474 public metadata also leave the inoculation protocol underspecified, including the inoculum form
475 (mycelia, sporangia, or zoospores), the developmental stage of the sampled leaves, and whether
476 the tissue was wounded before inoculation, so all conclusions are additionally conditional on the
477 depositors' experimental design. All quantification depends on a single reference assembly, and
478 although uniform and observed mapping checks removed the clearest failure modes, residual
479 reference bias cannot be excluded. The primary external cohort is deliberately heterogeneous
480 with respect to the discovery cohort, which makes support conservative and non-support
481 ambiguous; no genuinely independent replication cohort (same tissue, same time, different
482 biological material) exists in the public record for this design. Unsupported external outcomes
483 additionally conflate limited power, context dependence, and true absence of effect. The
484 simulations quantify the first component: conditional 80% minimum detectable effects were 2.44
485 log₂FC (5.4-fold) for genome-wide discovery and 2.11 log₂FC (4.3-fold) for candidate-family
486 external evaluation, with still poorer performance among low-expression genes. These
487 conditional curves do not identify which unsupported genes are context-specific or truly null.
488 And the entire study is computational: no candidate, including the Tier B set, has been perturbed
489 experimentally. The natural next steps are quantitative PCR of the Tier B genes in independently
490 infected material of both discovery cultivars, functional tests of the alpha-xylosidase and DIR1-
491 like candidates, and, for any regulatory hypothesis, direct assays such as promoter-reporter or
492 chromatin-binding experiments rather than further motif counting.

4. Conclusions

By carrying a single pathosystem from exploratory candidate discovery through a prospectively registered genome-wide discovery framework with prespecified cross-context external evaluation, this study establishes which claims about cultivar-dependent responses of lychee to *P. litchii* the current public data can and cannot support. It can support a set of 206 QC-retained genome-wide statistical interaction candidates, of which 19 are significant under both statistical frameworks, a stable core of 16 internally robust genes, 12 Tier B candidates whose annotations include cell-wall, sugar-metabolism, lipid-signaling, and RNA-metabolism functions, and 2 genes with cross-context support in an independent tissue and cultivar contrast. It cannot support the original 18 exploratory candidates as validated markers, any pathway- or signature-level claim of context-free transport, or any regulatory-element claim under matched-background testing. All frozen results, including nulls, contradictions, and retirements, are released with the analysis code and checksummed source data, so that the next stage of experimental validation can begin from an honest baseline.

5. Materials and Methods

5.1 Study design, registration, and evidence vocabulary

The confirmatory stage was governed by a time-stamped protocol, with amendments logged, distributed with the accompanying repository. The registration chronology is as follows. The exploratory stage was completed first, so the discovery dataset had already been analyzed, under different models, before registration. Protocol version 1.0.0 was frozen on 18 August 2026, with its SHA-256 hash (ef668486...) recorded alongside the protocol text, before any confirmatory expression outcome was computed. The amendment log records 20 outcome-blind amendments (clarifications, implementation corrections, and optimizations, all time-stamped before any

discovery or external expression outcome was viewed) and 3 post-outcome amendments confined to technical quality control and manifest verification, each marked as such. The primary external study had been published previously [5], and its headline biological conclusions were known to the analyst before registration; dataset roles and all thresholds were nevertheless fixed before any external count matrix was processed. Because all cohorts are public, the staged external unlock is a workflow-enforced computational safeguard, not blinding in the experimental sense. The protocol fixed dataset roles (Table 1), statistical models, discovery and evaluation thresholds, conditional gates, empty-set behavior, and the vocabulary used throughout this paper: *frozen* denotes a computed result fixed at a recorded checkpoint and never modified afterward; *discovery* refers only to the registered genome-wide interaction analysis in PRJNA830488; *robustness* refers only to predeclared sensitivity analyses within that cohort; *cross-context support* denotes a passing frozen external test in a cohort that differs in tissue, comparator, or design, and is never described as replication; and *orthogonal* evidence classes (annotation, motifs, small RNA, literature) are reported separately from expression support and cannot substitute for it. No external outcome could modify frozen discovery membership, signature weights, or thresholds.

5.2 Exploratory-stage analyses (Stage 1)

The exploratory stage analyzed the 12 GSE201243 libraries with FastQC v0.11.9 quality control, Trimmomatic v0.39 trimming, TopHat v2.0.12/Bowtie2 v2.2.3 alignment to the SCAU_Lch_v2.0 assembly (GCA_019925255.1), HTSeq v2.0.2 union-mode counting, and DESeq2 v1.42.0 per-cultivar Wald contrasts (design ~ treatment) with Benjamini–Hochberg adjustment at $P < 0.05$ [10–15]. Candidate prioritization combined significance, fold-change magnitude, and annotation plausibility; structural, conserved-motif, and promoter analyses used TBtools-II v2.096, MEME v5.5.5, and PlantCARE [16–18]; and the promoter-background

comparison tested observed *cis*-element counts in the 18 candidate promoters against 100,000 randomized sets of eighteen 2-kb sequences with lychee-like GC content, scoring exact motifs and reverse complements, with one-sided empirical *P*-values adjusted by Benjamini–Hochberg [19]. The selected-gene interaction contrast was computed as $\log_2\text{FC}(\text{Yurong1, infected versus mock}) - \log_2\text{FC}(\text{Guiwei, infected versus mock})$ over the prioritized genes and was explicitly labeled a prioritization effect size rather than a genome-wide test.

5.3 Data provenance and eligibility

National Center for Biotechnology Information and Gene Expression Omnibus metadata for all five cohorts were reconciled into a biological-unit registry (Supplementary S1). The original repository preserved accession decisions but not the historical query log, so search coverage was reconstructed retrospectively on 27 August 2026 and is not represented as prospective provenance. A broad NCBI GEO series query returned six records and a corresponding BioProject query returned seven; a narrow SRA cross-check returned 42 experiment-level records. After collapsing the GSE222652 super-series, excluding PRJNA1157370 because it used *Colletotrichum gloeosporioides*, and excluding PRJNA268587/GSE63658 because it studied fruit senescence/storage without *P. litchii* inoculation, the five prespecified cohorts remained (Supplementary S16a–b). Because source-tree, pooling, harvest, and extraction independence were not reported for the discovery libraries, all inferential *q*-values are conditional on deposited-library independence. The primary model required at least three included libraries in every cultivar-treatment cell.

5.4 Confirmatory preprocessing and quality gates

Host (SCAU_Lch_v2.0) and pathogen references were combined with prefixed contig names to absorb pathogen-derived reads. Paired reads were checked with FastQC, trimmed with fastp

[20], aligned with STAR two-pass mapping [21], counted at the gene level with featureCounts [22], and independently quantified with Salmon [23]. Libraries were excluded if fewer than 10 million read pairs survived trimming or fewer than 40% of reads aligned uniquely; all 12 discovery libraries passed. GenMap-based exon-uniqueness scores [24] and candidate-level observed-mapping checks were fixed before discovery and used for the uniform mapping and observed mapping-sensitivity gates.

5.5 Genome-wide discovery model

DESeq2 [15] fitted gene counts with the design $\sim \text{cultivar} + \text{treatment} + \text{cultivar}:\text{treatment}$ (reference levels Guiwei, mock). The primary coefficient, `cultivarYurong1:treatmentinfected`, estimates the difference in infected-minus-mock response between Yurong1 and Guiwei. Genes required counts of at least 10 in at least three libraries; statistical candidates required genome-wide Benjamini–Hochberg $q < 0.05$ and $|\log_2\text{FC}| \geq \log_2(1.5)$ [19]. The registered rule tests the point null of zero interaction and applies the effect threshold to the observed estimate. Post hoc sensitivity used DESeq2 with `lfcThreshold = log2(1.5)` and `altHypothesis = greaterAbs`, plus `apeglm` false-sign-or-small s -values for the same threshold; neither analysis modified frozen membership. To align the legacy audit with its original estimand, separate post hoc DESeq2 models ($\sim \text{treatment}$) were fitted within Guiwei and Yurong1, using the same count matrix, 19,445-gene universe, and genome-wide Benjamini–Hochberg adjustment within each cultivar. Unshrunk maximum-likelihood estimates are reported throughout; `apeglm`-shrunk interaction estimates [25] were used only to define the signed signature weights over the 206 frozen genes. Mappability and gene-model metrics were computed for all 19,445 tested genes; as a sensitivity check, recomputing the Benjamini–Hochberg adjustment over only the 13,602 QC-eligible genes rediscovered all 206 frozen candidates plus 12 additional genes, confirming that applying quality control after significance testing was conservative here.

5.6 Pathway and transcript-usage discovery

Plant Reactome pathway memberships [26] were transported to lychee through one-to-one reciprocal-best DIAMOND protein matches to *Oryza* [27], and pathway discovery applied the frozen genome-wide interaction statistic to this fixed mapping with camera, roast, and fgsea testing plus leading-edge-deletion and expression/length-matched null gates for robustness [28–30]. Differential transcript usage used Salmon transcript abundances analyzed with DRIMSeq and DEXSeq under stageR stage-wise error control [31–33]; the conditional inferential gate for interpretable two-level error control was registered in advance, and no substitute analysis was permitted had it failed.

5.7 Internal robustness gates

Each frozen gene was retested under Salmon-based gene-level quantification, edgeR quasi-likelihood interaction testing [34], a counts-per-million expression filter, and all leave-one-library-out refits, with predeclared criteria combined conjunctively; no additive scoring was used. The Salmon gate required direction agreement and an absolute effect difference of at most 0.5 log₂ units; the edgeR analysis used the identical design matrix and interaction coefficient with TMM normalization, robust dispersion estimation (estimateDisp and glmQLFit with robust settings), and glmQLFTest, with Benjamini–Hochberg adjustment computed genome-wide over all 19,445 genes and a gate of genome-wide $q < 0.10$ plus direction agreement; the expression-filter gate refit the primary model after a counts-per-million filter; and the leave-one-library-out gate required direction agreement in all 12 refits and $q < 0.10$ in at least 10 of 12. The observed mapping-sensitivity gate removed candidates whose effect depended on ambiguously mapped reads. Frozen pathways required agreement across camera, roast, leading-edge deletion, and the matched-null comparison.

5.8 External frozen evaluation

PRJNA450886 was analyzed with the design $\sim \text{cultivar} \times \text{treatment} \times \text{time}$, with time encoded as a categorical factor and 24 h as the reference level, so that the cultivar-by-treatment coefficient estimates the 24-h interaction and the three-way terms give the 6-h and 48-h contrasts; the primary estimand was the resistant-minus-susceptible cultivar-by-infection interaction at 24 h, with 6-h and 48-h contrasts secondary. Frozen candidate tests used Benjamini–Hochberg adjustment across the full candidate-by-contrast family, an absolute effect threshold of $\log_2(1.5)$, confidence intervals excluding zero, and discovery-direction agreement; opposite-direction passes were recorded as contradictory. PRJNA922966 was analyzed with $\sim \text{tissue} + \text{treatment} + \text{tissue:treatment}$ for generic-transfer tests, and PRJNA1090613 outcomes were confined to exploratory status by protocol. The frozen signed signature was evaluated as a library-level score with the same family-wise adjustment.

5.9 Annotation, motif, small-RNA, and literature classes

Representative proteins were chosen by longest coding sequence with lexicographic tie-breaking. Reviewed Swiss-Prot (release 2026_02) DIAMOND matches [35], targeted InterPro/Pfam domain assignments [36], and one-to-one *Oryza* orthology formed three separate annotation classes; the high-confidence label required at least two classes, at least 70% supported coverage, and no architecture conflict. Promoter-motif analysis tested 927 JASPAR plant position-weight matrices [37] in strand-aware 1-kb and 2-kb promoter windows against 100 expression- and GC-matched background sets, with AME as the primary test, FIMO as sensitivity [38,39], robustness requiring passes in at least 80 of 100 backgrounds in both windows, and cognate transcription-factor expression required for attribution. The PmiREN bulk archive [40] was audited before small-RNA analysis; because it contained no exact *Litchi* entries, the class was frozen as not testable and no study-derived reference was substituted. Published studies reusing any analyzed

accession were registered as prior interpretation. For the post hoc controlled PlantCARE comparison, published observed totals across the 18 legacy promoters were fixed and tested against (i) 100,000 simulated sets of eighteen 2-kb sequences at 34% GC and (ii) 100 genomic sets of eighteen promoters matched by mean-expression and promoter-GC quintiles with the same seeded minimum-distance rule used by the registered motif pipeline. The six reported motifs and one standard TCA variant, together with reverse complements, were counted and one-sided empirical *P*-values were adjusted across seven classes. The missing archived observed-motif TSV and canonical-promoter recount discrepancies were retained as explicit provenance limitations (Supplementary S15).

5.10 Simulation-based power analysis

Negative-binomial parametric simulations used DESeq2 maximum-a-posteriori gene dispersions and fitted sample means from the discovery counts and from the 24-h subset of PRJNA450886. Fitted cultivar and treatment main effects and size factors were retained, background-gene interaction coefficients were set to zero, and equal numbers of positive and negative effects were injected at 12 absolute log₂FC values (0, 0.5, 0.75, 1, 1.25, 1.5, 1.75, 2, 2.5, 3, 3.5, 4). Each value had 100 simulations. Discovery injected 262 targets and applied Benjamini–Hochberg correction over 19,445 genes plus the registered effect filter. External simulation injected 2 targets among the 177 frozen genes passing the simulation count and dispersion validity filters and additionally required the prespecified sign and a 95% interval excluding zero; the observed external table contained 184 measurable genes, seven of which were not fit-eligible for the simulation. Dispersions were treated as known during simulated Wald refits, making the curves conditional parametric power rather than a claim about population-level biological replication. Eighty-percent minimum detectable effects were linearly interpolated between grid points, with expression-quartile curves and Wilson intervals reported in Figure S2.

5.11 Tier rules, software versions, and reproducibility

Tier A required discovery, complete internal robustness, primary cross-context support, clean mapping and annotation, and at least one attributable orthogonal class; Tier B required discovery and complete internal robustness with partial or non-testable primary external evidence; Tier C was reserved for a robust, externally supported pathway. Contradictory external direction or mapping/annotation failure retired an entity, and an empty Tier A was an allowed outcome. All stages ran under Snakemake 9.25.2 workflows [41] enforcing staged discovery, a controlled external unlock, verified cleanup, and SHA-256 manifests over inputs and outputs. Confirmatory executables were FastQC 0.12.1, fastp 1.3.6, STAR 2.7.10b, featureCounts/Subread 2.1.1, Salmon 2.5.1, MEME Suite/AME/FIMO 5.5.9, and DIAMOND 2.2.5; GenMap 1.3.0 is recoverable from the pinned specification although its executable is absent from the reconstructed environment. Statistical work used R 4.3.3/Bioconductor 3.18 with DESeq2 1.42.0, apegglm 1.24.0, edgeR 4.0.16, DRIMSeq 1.30.0, DEXSeq 1.48.0, stageR 1.24.0, and BiocParallel 1.36.0 (Supplementary S17 records the evidence status of every version). The resolved environment, synthetic fixtures, per-figure source data, and amendment log are distributed with the repository. All six main figures are regenerated from frozen result tables by a script that recomputes and asserts every headline count.

Acknowledgments

The author has no acknowledgments to declare.

Funding

This work received no external funding.

Author Contributions

E.Z. conceived the study, performed all analyses, and wrote the manuscript. The author approved the submitted version and is accountable for the work.

Conflict of Interest

The author declares no conflicts of interest.

Data Availability

All analyzed sequencing data are public under PRJNA830488/GSE201243, PRJNA450886, PRJNA922966/GSE222651, PRJNA922965/GSE222650, and PRJNA1090613/GSE262200. Supplementary Tables S1-S18, Figures S1-S3, and tab-separated figure source data are archived under CC BY 4.0 at <https://doi.org/10.5281/zenodo.22240717>.

Accession Numbers

Sequence data analyzed in this article can be found in the NCBI SRA and GEO repositories under PRJNA830488 (GSE201243), PRJNA450886, PRJNA922966 (GSE222651), PRJNA922965 (GSE222650), and PRJNA1090613 (GSE262200). Supplemental data are permanently archived at <https://doi.org/10.5281/zenodo.22240717>.

Supplemental Data

The following items are archived in the supplement at <https://doi.org/10.5281/zenodo.22240717>:

- Table S1. Biological-unit registry.
- Table S2. Per-library quality control.
- Table S3. Genome-wide discovery statistics.
- Table S4. Robustness file manifest.
- Table S5. All external frozen tests.

- 701 • Table S6. Pathway and signature tests.
- 702 • Table S7. Conditional differential-transcript-usage results.
- 703 • Table S8. Candidate annotation and orthology evidence.
- 704 • Table S9. Small-RNA reference-gate outcome.
- 705 • Table S10. Motif-background tests and sensitivity results.
- 706 • Table S11. Accession-aware evidence registry.
- 707 • Table S12. Scripts, environments, and command inventory.
- 708 • Table S13. Protocol amendment and deviation log.
- 709 • Table S14. Legacy within-cultivar audit.
- 710 • Table S15. Controlled promoter-motif background comparison.
- 711 • Table S16a. Reconstructed dataset-search queries.
- 712 • Table S16b. Dataset eligibility decisions.
- 713 • Table S17. Exact software versions.
- 714 • Table S18. Power-simulation minimum detectable effects.
- 715 • Figure S1. Replicate-level normalized counts.
- 716 • Figure S2. Conditional parametric power curves.
- 717 • Figure S3. Exploratory signed-signature estimate.
- 718 • Data S1. Tab-separated source data for all main and supplementary analytical figures.

References

1. Morton JF. Lychee. In: Fruits of Warm Climates. Winterville: Creative Resource Systems; 1987. p. 249–259.
2. Sun J, Gao Z, Zhang X, Zou X, Cao L, Wang J. Transcriptome analysis of *Phytophthora litchii* reveals pathogenicity arsenals and confirms taxonomic status. PLoS ONE. 2017;12:e0178245. doi:10.1371/journal.pone.0178245.
3. Yi C, Jiang Y, Shi J, Qu H, Xue S, Duan X, et al. ATP-regulation of antioxidant properties and phenolics in litchi fruit during browning and pathogen infection process. Food Chemistry. 2010;118:42–47. doi:10.1016/j.foodchem.2009.04.074.
4. Hu G, Feng J, Xiang X, Wang J, Salojärvi J, Liu C, et al. Two divergent haplotypes from a highly heterozygous lychee genome suggest independent domestication events for early and late-maturing cultivars. Nature Genetics. 2022;54:73–83. doi:10.1038/s41588-021-00971-3.
5. Sun J, Cao L, Li H, Wang G, Wang S, Li F, et al. Early responses given distinct tactics to infection of *Peronophythora litchii* in susceptible and resistant litchi cultivar. Scientific Reports. 2019;9:2810. doi:10.1038/s41598-019-39100-w.
6. Liu H, Yan Q, Jiang Y, Shi F, Chen J, Cai C, et al. Identification of LcCDPKs and analysis of their expression patterns in response to downy mildew stresses in lychee. Journal of Fruit Science. 2023;40:442–456. doi:10.13925/j.cnki.gsxb.20220307.
7. Li P, Li W, Zhou X, Situ J, Xie L, Xi P, et al. *Peronophythora litchii* RXLR effector PlAvh202 destabilizes a host ethylene biosynthesis enzyme. Plant Physiology. 2023;193:756–774. doi:10.1093/plphys/kiad311.
8. National Center for Biotechnology Information. GEO accession GSE201243. <https://www.ncbi.nlm.nih.gov/geo/query/acc.cgi?acc=GSE201243>.
9. Berka M, Kopecká R, Berková V, Brzobohatý B, Černý M. Regulation of heat shock proteins 70 and their role in plant immunity. Journal of Experimental Botany. 2022;73:1894–1909. doi:10.1093/jxb/erab549.
10. Andrews S. FastQC: a quality control tool for high throughput sequence data. Babraham Bioinformatics; 2010.
11. Bolger AM, Lohse M, Usadel B. Trimmomatic: a flexible trimmer for Illumina sequence data. Bioinformatics. 2014;30:2114–2120. doi:10.1093/bioinformatics/btu170.
12. Kim D, Pertea G, Trapnell C, Pimentel H, Kelley R, Salzberg SL. TopHat2: accurate alignment of transcriptomes in the presence of insertions, deletions and gene fusions. Genome Biology. 2013;14:R36. doi:10.1186/gb-2013-14-4-r36.
13. Anders S, Pyl PT, Huber W. HTSeq—a Python framework to work with high-throughput sequencing data. Bioinformatics. 2015;31:166–169. doi:10.1093/bioinformatics/btu638.
14. Genome Warehouse. *Litchi chinensis* SCAU_Lch_v2.0 assembly (GCA_019925255.1). National Genomics Data Center, China National Center for Bioinformation.

- 756 15. Love MI, Huber W, Anders S. Moderated estimation of fold change and dispersion for RNA-
757 seq data with DESeq2. *Genome Biology*. 2014;15:550. doi:10.1186/s13059-014-0550-8.
- 758 16. Chen C, Chen H, Zhang Y, Thomas HR, Frank MH, He Y, et al. TBtools: an integrative
759 toolkit developed for interactive analyses of big biological data. *Molecular Plant*.
760 2020;13:1194–1202. doi:10.1016/j.molp.2020.06.009.
- 761 17. Bailey TL, Elkan C. Fitting a mixture model by expectation maximization to discover motifs
762 in biopolymers. In: *Proceedings of the Second International Conference on Intelligent*
763 *Systems for Molecular Biology*. Menlo Park: AAAI Press; 1994. p. 28–36.
- 764 18. Lescot M, Déhais P, Thijs G, Marchal K, Moreau Y, Van de Peer Y, et al. PlantCARE, a
765 database of plant cis-acting regulatory elements and a portal to tools for in silico analysis of
766 promoter sequences. *Nucleic Acids Research*. 2002;30:325–327. doi:10.1093/nar/30.1.325.
- 767 19. Benjamini Y, Hochberg Y. Controlling the false discovery rate: a practical and powerful
768 approach to multiple testing. *Journal of the Royal Statistical Society, Series B*. 1995;57:289–
769 300.
- 770 20. Chen S, Zhou Y, Chen Y, Gu J. fastp: an ultra-fast all-in-one FASTQ preprocessor.
771 *Bioinformatics*. 2018;34:i884–i890. doi:10.1093/bioinformatics/bty560.
- 772 21. Dobin A, Davis CA, Schlesinger F, Drenkow J, Zaleski C, Jha S, et al. STAR: ultrafast
773 universal RNA-seq aligner. *Bioinformatics*. 2013;29:15–21.
774 doi:10.1093/bioinformatics/bts635.
- 775 22. Liao Y, Smyth GK, Shi W. featureCounts: an efficient general purpose program for assigning
776 sequence reads to genomic features. *Bioinformatics*. 2014;30:923–930.
777 doi:10.1093/bioinformatics/btt656.
- 778 23. Patro R, Duggal G, Love MI, Irizarry RA, Kingsford C. Salmon provides fast and bias-aware
779 quantification of transcript expression. *Nature Methods*. 2017;14:417–419.
780 doi:10.1038/nmeth.4197.
- 781 24. Pockrandt C, Alzamel M, Iliopoulos CS, Reinert K. GenMap: ultra-fast computation of
782 genome mappability. *Bioinformatics*. 2020;36:3687–3692.
783 doi:10.1093/bioinformatics/btaa222.
- 784 25. Zhu A, Ibrahim JG, Love MI. Heavy-tailed prior distributions for sequence count data:
785 removing the noise and preserving large differences. *Bioinformatics*. 2019;35:2084–2092.
786 doi:10.1093/bioinformatics/bty895.
- 787 26. Naithani S, Gupta P, Preece J, D'Eustachio P, Elser J, Garg P, et al. Plant Reactome
788 knowledgebase: empowering plant pathway exploration and OMICS data analysis. *Nucleic*
789 *Acids Research*. 2024;52:D1538–D1547. doi:10.1093/nar/gkad1052.
- 790 27. Buchfink B, Xie C, Huson DH. Fast and sensitive protein alignment using DIAMOND.
791 *Nature Methods*. 2015;12:59–60. doi:10.1038/nmeth.3176.
- 792 28. Wu D, Smyth GK. Camera: a competitive gene set test accounting for inter-gene correlation.
793 *Nucleic Acids Research*. 2012;40:e133. doi:10.1093/nar/gks461.

- 794 29. Wu D, Lim E, Vaillant F, Asselin-Labat ML, Visvader JE, Smyth GK. ROAST: rotation
795 gene set tests for complex microarray experiments. *Bioinformatics*. 2010;26:2176–2182.
796 doi:10.1093/bioinformatics/btq401.
- 797 30. Korotkevich G, Sukhov V, Budin N, Shpak B, Artyomov MN, Sergushichev A. Fast gene set
798 enrichment analysis. *bioRxiv*. 2021:060012. doi:10.1101/060012.
- 799 31. Nowicka M, Robinson MD. DRIMSeq: a Dirichlet-multinomial framework for multivariate
800 count outcomes in genomics. *F1000Research*. 2016;5:1356.
801 doi:10.12688/f1000research.8900.2.
- 802 32. Anders S, Reyes A, Huber W. Detecting differential usage of exons from RNA-seq data.
803 *Genome Research*. 2012;22:2008–2017. doi:10.1101/gr.133744.111.
- 804 33. Van den Berge K, Sonesson C, Robinson MD, Clement L. stageR: a general stage-wise
805 method for controlling the gene-level false discovery rate in differential expression and
806 differential transcript usage. *Genome Biology*. 2017;18:151. doi:10.1186/s13059-017-1277-
807 0.
- 808 34. Robinson MD, McCarthy DJ, Smyth GK. edgeR: a Bioconductor package for differential
809 expression analysis of digital gene expression data. *Bioinformatics*. 2010;26:139–140.
810 doi:10.1093/bioinformatics/btp616.
- 811 35. The UniProt Consortium. UniProt: the Universal Protein Knowledgebase in 2023. *Nucleic
812 Acids Research*. 2023;51:D523–D531. doi:10.1093/nar/gkac1052.
- 813 36. Jones P, Binns D, Chang HY, Fraser M, Li W, McAnulla C, et al. InterProScan 5: genome-
814 scale protein function classification. *Bioinformatics*. 2014;30:1236–1240.
815 doi:10.1093/bioinformatics/btu031.
- 816 37. Rauluseviciute I, Riudavets-Puig R, Blanc-Mathieu R, Castro-Mondragon JA, Ferenc K,
817 Kumar V, et al. JASPAR 2024: 20th anniversary of the open-access database of transcription
818 factor binding profiles. *Nucleic Acids Research*. 2024;52:D174–D182.
819 doi:10.1093/nar/gkad1059.
- 820 38. McLeay RC, Bailey TL. Motif enrichment analysis: a unified framework and an evaluation
821 on ChIP data. *BMC Bioinformatics*. 2010;11:165. doi:10.1186/1471-2105-11-165.
- 822 39. Grant CE, Bailey TL, Noble WS. FIMO: scanning for occurrences of a given motif.
823 *Bioinformatics*. 2011;27:1017–1018. doi:10.1093/bioinformatics/btr064.
- 824 40. Guo Z, Kuang Z, Wang Y, Zhao Y, Tao Y, Cheng C, et al. PmiREN: a comprehensive
825 encyclopedia of plant miRNAs. *Nucleic Acids Research*. 2020;48:D1114–D1121.
826 doi:10.1093/nar/gkz894.
- 827 41. Mölder F, Jablonski KP, Letcher B, Hall MB, Tomkins-Tinch CH, Sochat V, et al.
828 Sustainable data analysis with Snakemake. *F1000Research*. 2021;10:33.
829 doi:10.12688/f1000research.29032.2.

830 Figure Legends

831 **Figure 1.** Study cohorts, discovery quality control, and expression structure. (A) The five public
832 cohorts with prespecified roles and deposited library counts. (B) Per-library technical quality in
833 the discovery cohort: fastp read survival, STAR unique-alignment rate, and Salmon mapping
834 rate; all 12 libraries passed the fixed inclusion gates. (C) Principal-component analysis of the
835 discovery libraries separates cultivar and infection status.

836 **Figure 2.** Genome-wide discovery and formal reassessment of the exploratory candidates. (A)
837 Cultivar-by-infection interaction effects across 19,445 expressed genes; 262 genes met the
838 statistical threshold and 206 were retained after uniform mapping and gene-model quality
839 control. (B) Genome-wide re-estimation of the 18 exploratory (legacy) candidates with 95%
840 confidence intervals; orange marks uniform-mappability failures. No legacy candidate met the
841 genome-wide interaction criterion. (C) Evidence attrition from tested genes to internally robust
842 candidates.


843 **Figure 3.** Internal robustness and the primary external evaluation. (A) Candidates passing each
844 predeclared robustness gate, the conjunction of the four analytic gates, and the additional
845 observed-mapping gate. (B) The internally robust set (16 genes) and the externally supported set
846 (2 genes) do not overlap. (C) Discovery interaction effects against the primary cross-context
847 interaction effects (Heiye response minus Guiwei response, pericarp, 24 h); supported,
848 contradictory, and internally robust genes are highlighted. The diagonal is a visual reference
849 only; equal effects across contexts are not expected.

850 **Figure 4.** Frozen pathway and signature evaluation. (A) Discovery normalized enrichment
851 scores for the six frozen pathways; only circadian rhythm passed all internal gates, and none was
852 supported in the primary external evaluation. (B) The frozen signed 206-gene signature across
853 the six non-exploratory external contrasts; filled points denote $q < 0.05$. The quarantined
854 PRJNA1090613 estimate is in Figure S3.


855 **Figure 5.** Conditional transcript usage. (A) Discovery gate and event counts. (B) External
856 follow-up of the 225 discovery events, separated by prespecified study role.

857 **Figure 6.** Orthogonal evidence and final tiers. (A) Annotation status of the 206 QC-retained
858 genes. (B) Registered promoter-motif robustness and the small-RNA reference gate. (C)
859 Deterministic final tiers over all 268 frozen entities.


A Cohort roles and deposited libraries



B Discovery technical QC




C Expression PCA separates cultivar and infection response





A Genome-wide interaction discovery and uniform gene QC

Page 41 of 48




B Legacy 18: formal genome-wide reassessment


C Evidence attrition




A Internal robustness gates



B No gene-level convergence




C Discovery versus primary cross-context effects





A Only the circadian pathway passed internal robustness

Page 43 of 48



B The signed discovery score changed with time and tissue






Stage-wise q values and signed usage-interaction differences

B External DTU follow-up remained separated by study role

Supported Contradictory Unsupported Not testable



No primary event was supported; exploratory outcomes remain non-confirmatory.



High-confidence gate 0 high-confidence annotations




Figure S1





Figure S1. Replicate-level normalized counts for the two cross-context-supported genes, twelve Tier B genes, and two legacy highlights. Points are individual deposited libraries; bars show cultivar-treatment medians.

Figure S2



Dashed line: 80% power. Dotted line: registered $|\log_2\text{FC}|$ threshold. Parametric power is conditional on fitted means, dispersions, deposited-library independence, and the simulated non-null prevalence.

Figure S2. Parametric detection probability for cultivar-by-infection effects under genome-wide discovery and candidate-family external adjustment. Curves show the overall result and mean-expression quartiles; the dashed line marks 80% power.


Figure S3

Figure S3. Quarantined PRJNA1090613 signed-signature estimate, retained as exploratory rather than confirmatory evidence.