

ROYAL SOCIETY OPEN SCIENCE

Cultivar-dependent transcriptional responses of lychee to *Peronophythora litchii*: a registered genome-wide analysis

Journal:	<i>Royal Society Open Science</i>
Manuscript ID	Draft
Article Type:	Research
Date Submitted by the Author:	n/a
Complete List of Authors:	Zhuang, Eric ; NYU Langone Health
Subject:	bioinformatics < BIOLOGY
Keywords:	Litchi chinensis, Peronophythora litchii, plant-pathogen interaction, RNA sequencing, cultivar-dependent response, preregistered analysis
Subject Category:	Biochemistry, Cellular and Molecular Biology
Note: The following files were submitted by the author for peer review, but cannot be converted to PDF. You must view these files (e.g. movies) online.	
S13_amendment_deviation_log.tsv S14_legacy_within_cultivar_audit.tsv S17_exact_tool_versions.tsv S18_power_simulation_mde.tsv S1_biological_unit_registry.tsv S2_per_library_QC.tsv S3_all_discovery_statistics.tsv S5_all_external_frozen_tests.tsv S6_pathway_signature_tests.tsv S7_DTU_results.tsv S8_annotation_orthology.tsv S9_small_RNA_results.tsv	



1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60

Statement 2026

Author-supplied statements

Relevant information will appear here if provided.

Ethics

Does your article include research that required ethical approval or permits?:
This article does not present research with ethical considerations

Statement (if applicable):
CUST_IF_YES_ETHICS :No data available.

Data

It is a condition of publication that data, code and materials supporting your paper are made publicly available. Does your paper present new data?:
Yes

Statement (if applicable):
Here is the link to the supplemental data for this paper:
<https://doi.org/10.5281/zenodo.22240717>

All analysed sequencing data are public under PRJNA830488/GSE201243, PRJNA450886, PRJNA922966/GSE222651, PRJNA922965/GSE222650 and PRJNA1090613/GSE262200. Supplementary Tables S1–S18, Figures S1–S3 and tab-separated figure source data are archived under CC BY 4.0 at <https://doi.org/10.5281/zenodo.22240717>. The analysis source code, Snakemake workflows, configuration files, tests and environment specifications are supplied within the electronic supplementary files. The data, code and supporting materials required to reproduce the reported results are therefore available to editors and reviewers with this submission.

Conflict of interest

I/We declare we have no competing interests

Statement (if applicable):
CUST_STATE_CONFLICT :No data available.

Use of AI

Please provide a statement of any use of AI technology in the preparation of the paper.

No, we have not used AI-assisted technologies in creating this article
CUST_IF_YES_DECLARATION_OF_AI_USE :No data available.

**Cultivar-dependent transcriptional responses of lychee to
Peronophythora litchii: a registered genome-wide analysis**

Eric Zhuang

NYU Langone Health, New York, NY, USA

Corresponding author: Eric Zhuang

Email: eric.zhuang@nyulangone.org

ORCID iD: 0009-0001-9050-0214

Article type: Research Article

Abstract

Litchi downy blight, caused by the oomycete *Peronophythora litchii*, is a major disease of lychee. Public RNA-sequencing cohorts offer a resource for identifying cultivar-dependent infection responses, but are readily over-interpreted when candidate selection and confirmation use the same data. We combined an exploratory reanalysis with a prospectively registered genome-wide discovery and external-evaluation workflow. A negative-binomial cultivar-by-infection analysis of 19,445 expressed genes yielded 262 statistical candidates; 206 passed mapping and gene-model quality control, 19 remained significant under an independent statistical framework and 16 passed all registered robustness checks. In a prespecified cross-context evaluation using an independent pericarp time course, two genes were supported and five showed significant opposite-direction effects; the internally robust and externally supported sets did not overlap beyond chance expectation. None of 18 exploratory candidates met the subsequent genome-wide interaction criterion. Deterministic evidence integration produced no highest-confidence candidates, retained 12 internally robust genes in a middle tier and retired 65 entities. These results show that prominent single-cohort candidates often fail formal interaction testing and cross-context evaluation. The openly archived workflow, complete results and explicit negative findings provide a reproducible baseline for targeted experimental validation.

Keywords: *Litchi chinensis*; *Peronophythora litchii*; plant-pathogen interaction; RNA sequencing; cultivar-dependent response; preregistered analysis

1. Introduction

Lychee (*Litchi chinensis* Sonn.) is a subtropical fruit crop of substantial economic importance in southern China, Southeast Asia, and a growing set of production regions worldwide [1]. Among its diseases, litchi downy blight, caused by the oomycete *Peronophythora litchii* (historically also

33 treated as *Phytophthora litchii* [2]), is particularly damaging: the pathogen attacks young leaves,
34 inflorescences, and fruit, produces water-soaked brown lesions, and drives severe postharvest
35 decay [2,3]. Because chemical control is costly and resistance breeding depends on
36 understanding how host genotypes differ in their response to infection, the molecular basis of
37 cultivar-level variation in the response to *P. litchii* is a question of both agronomic and biological
38 interest.

39 The resources needed to address this question computationally have accumulated rapidly. High-
40 quality genome assemblies are available for lychee [4], and several groups have profiled the
41 transcriptional response to *P. litchii* infection. Sun et al. compared a susceptible and a resistant
42 cultivar across a pericarp infection time course and reported that the two genotypes deploy
43 distinct early response programs [5]. Liu et al. described calcium-dependent protein kinase
44 (LcCDPK) family members induced by *P. litchii* inoculation [6], and Li et al. showed that the
45 pathogen RXLR effector PIAvh202 promotes virulence by destabilizing a host ethylene-
46 biosynthesis enzyme [7]. A pathogen-side transcriptome catalogued the *P. litchii* pathogenicity
47 arsenal [2]. In parallel, public repositories now hold RNA-seq series covering resistant and
48 susceptible cultivars, leaf and fruit tissues, messenger and small RNA, and multiple time points.
49 These series offer an opportunity, in principle, to distinguish genome-wide cultivar-dependent
50 infection responses from responses common to all genotypes.

51 Realizing that opportunity is statistically demanding. The quantity of interest is not differential
52 expression within a cultivar but the cultivar-by-infection interaction: the difference between
53 infection responses across genotypes. Interaction effects are estimated with less precision than
54 main effects, the deposited designs are small (typically three libraries per cultivar-treatment
55 cell), and a genome-wide search requires an explicit multiplicity family. Analyses that instead

1
2
3 56 select a handful of strongly responding genes and then characterise them in the same cohort
4
5 57 conflate selection with confirmation: large observed fold changes in a small experiment are a
6
7 58 mixture of true signal and noise, and features chosen for extremity will regress toward null in
8
9 59 any independent assessment. These concerns apply directly to the present pathosystem, in which
10
11 60 most published candidate genes derive from single-cohort analyses without external evaluation.
12
13
14 61 The first stage of the present work followed the same pattern. In an exploratory reanalysis of the
15
16 62 GSE201243 leaf series [8] (Guiwei and Yurong1, mock versus inoculated, 24 h post-
17
18 63 inoculation), we identified 17 significant infection-responsive genes in Guiwei and 117 in
19
20 64 Yurong1, prioritised 18 candidates with defence-plausible annotations (among them jacalin-like
21
22 65 lectin, pathogenesis-related and thaumatin/osmotin-like proteins, a heat-shock protein, lipid-
23
24 66 transfer proteins, and calcium-binding proteins), and found that a set of stress- and hormone-
25
26 67 associated *cis*-regulatory elements was enriched in their promoters relative to a randomised
27
28 68 background. That analysis was framed as hypothesis-generating, and critical review of it
29
30 69 converged on three substantive criticisms: the cultivar-by-treatment comparison was an effect-
31
32 70 size contrast over pre-selected genes rather than a genome-wide interaction test with false-
33
34 71 discovery-rate control; no independent dataset was used for validation; and promoter-element
35
36 72 counting, even against a randomised background, is weak evidence of regulatory function. We
37
38 73 regard all three criticisms as correct, and they motivated the confirmatory stage reported here.
39
40 74 This paper therefore presents the complete arc of the project in one place. We prospectively
41
42 75 registered a discovery and external-evaluation protocol that fixed dataset roles, models,
43
44 76 thresholds, empty-result behaviour, and the vocabulary used to describe evidence before any
45
46 77 external outcome was opened. Because every cohort is public, this staged unlock is a
47
48 78 computational safeguard enforced by the workflow rather than blinding in the experimental
49
50
51
52
53
54
55
56
57
58
59
60

sense, and we document the registration chronology in the Methods. Within the prespecified discovery cohort we fitted a genome-wide negative-binomial interaction model, applied uniform mapping and gene-model quality control, and subjected the resulting candidates to a battery of internal robustness checks spanning independent quantification, an alternative statistical framework, filter sensitivity, and leave-one-library-out refits. We then evaluated the frozen candidates (throughout, *frozen* denotes a computed result fixed at a recorded checkpoint and not modified afterward), a frozen pathway set, and a frozen signed expression signature in independent public cohorts whose roles (primary cross-context evaluation, generic transfer, orthogonal modality, or exploratory) were fixed in advance. Orthogonal evidence classes (sequence annotation, promoter motifs, small-RNA coherence, and accession-aware literature) were kept separate from expression support, and a deterministic tier system integrated all layers while explicitly permitting an empty top tier. Finally, we audited the 18 exploratory candidates from the first stage against the genome-wide analysis. The result is a candidate resource for lychee downy blight whose every claim carries its evidence class, together with a case study in how much, and how little, of a plausible exploratory analysis survives formal statistical scrutiny.

2. Results

2.1 A two-stage design with prespecified dataset roles

The study comprises an exploratory stage completed first and a confirmatory stage whose protocol was registered before its analyses began (Figure 1A; Table 1). Five public cohorts enter the confirmatory stage in fixed, non-interchangeable roles. PRJNA830488 (GSE201243), the leaf series for Guiwei and Yurong1 at 24 h, is the sole discovery cohort; it is also the dataset analysed in the exploratory stage. PRJNA450886, a pericarp time course contrasting the

1
2
3 101 susceptible cultivar Guiwei with the resistant cultivar Heiye at 6, 24, and 48 h [5], serves as the
4
5 102 primary external evaluation. Because it differs from the discovery cohort in tissue, resistant
6
7 103 comparator, and time structure, we describe positive outcomes in it as *cross-context support*
8
9
10 104 rather than replication. PRJNA922966 (Feizixiao leaf and fruit, mock versus infected) tests
11
12 105 generic infection-response transfer but cannot address the cultivar interaction; PRJNA922965
13
14 106 provides small-RNA libraries as an orthogonal modality; and PRJNA1090613 (Guiwei versus
15
16 107 SFZ leaf) is exploratory only, because its time and resistance metadata could not be resolved.
17
18 108 Under the registered firewall, no external outcome could alter the frozen discovery membership,
19
20 109 signature weights, or thresholds, and the protocol explicitly allowed the top evidence tier to be
21
22
23
24 110 empty.

25
26
27 111 [Insert Figure 1 near here]

28
29 112 **Table 1. Locked dataset roles and eligibility.**

Accession	Design	Locked role	Evidence class
PRJNA830488 (GSE201243)	Guiwei/Yurong1 leaf; mock/infected; 24 h; 3 libraries per cell	Discovery	Conditional inferential
PRJNA450886	Guiwei/Heiye pericarp; mock/infected; 6, 24, 48 h; 3 libraries per cell	Primary external evaluation	Cross-context
PRJNA922966 (GSE222651)	Feizixiao leaf/fruit; mock/infected; 24 h; 3 libraries per cell	Generic infection transfer	Cross-context
PRJNA922965 (GSE222650)	Feizixiao leaf/fruit small RNA; mock/infected; 24 h; 3 libraries per cell	Orthogonal modality	Orthogonal
PRJNA1090613 (GSE262200)	Guiwei/SFZ leaf; mock/infected; 3 libraries per cell	Exploratory only	Exploratory

30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46 113
47
48 114 **2.2 Stage 1: exploratory candidate discovery in GSE201243**

49
50
51 115 The exploratory stage analysed the 12 discovery libraries with per-cultivar differential-
52
53 116 expression contrasts (infected versus mock) and identified 17 significant genes in Guiwei (15 up-
54
55
56 117 regulated, 2 down-regulated) and 117 in Yurong1 (39 up-regulated, 78 down-regulated) at

Benjamini–Hochberg adjusted $P < 0.05$, suggesting a broader transcriptional response in Yurong1. From these lists, 18 genes were prioritised by combining statistical significance in at least one cultivar, large infection-associated fold changes, and annotations plausibly related to defence: among them a jacalin-like lectin (LITCHI028104), a pathogenesis-related protein (LITCHI011394), osmotin/thaumatin-like proteins (LITCHI009301, LITCHI002793), an HSP70-family chaperone (LITCHI012388), lipid-transfer proteins (LITCHI019183, LITCHI018043), an EF-hand calcium-binding protein (LITCHI019519), and a phytosulfokine precursor (LITCHI007102). A selected-gene contrast of infection responses between cultivars nominated LITCHI019519, LITCHI001510, LITCHI028401, and LITCHI017676 as Guiwei-skewed and LITCHI028104 and LITCHI019183 as Yurong1-skewed candidates.

The exploratory stage also examined 2-kb upstream promoter regions of the 18 candidates. Against 100,000 randomised GC-matched promoter sets, four *cis*-element classes appeared enriched: anaerobic-response elements (ARE; 46 observed versus 24.6 expected, $q = 0.0003$), methyl-jasmonate-responsive motifs (54 versus 38.4, $q = 0.018$), TCA elements (9 versus 0.19, $q = 0.00003$), and TC-rich repeats (10 versus 0.15, $q = 0.00003$). ABRE and MYB-binding-site motifs, in contrast, were frequent but not enriched. All of these results were obtained by selecting and characterising genes within a single cohort; they carry no genome-wide multiplicity control and no external evaluation, and we treat them below strictly as hypotheses submitted to the confirmatory stage.

2.3 Discovery-cohort quality under the registered protocol

The confirmatory pipeline combined a joint host and pathogen reference, fastp trimming, STAR two-pass alignment, featureCounts gene counting, and independent Salmon quantification. Under this pipeline, all 12 discovery libraries passed the fixed technical gates of at least 10 million

1
2
3 141 surviving read pairs and at least 40% uniquely aligned reads (Figure 1B). Read survival after
4
5 142 trimming ranged from 99.56% to 99.66%, unique-alignment rates from 84.44% to 86.22%, and
6
7
8 143 Salmon mapping rates from 80.77% to 82.56%. Principal-component analysis of the transformed
9
10 144 counts separated libraries by cultivar on the first component (54.2% of variance) and by infection
11
12 145 status on the second (16.1%), with no library isolated from its replicate group (Figure 1C).
13
14 146 Because source-tree, pooling, and extraction independence are not reported for the deposited
15
16
17 147 libraries, all inferential statements below are conditional on deposited-library independence, a
18
19 148 limitation we return to in the Discussion.

21
22 149 **2.4 Genome-wide cultivar-by-infection discovery**

23
24
25 150 The primary discovery model fitted all expressed genes with the design ~ cultivar + treatment +
26
27 151 cultivar:treatment, with Guiwei and mock as reference levels, so that the interaction coefficient
28
29
30 152 estimates how much the infected-minus-mock response in Yurong1 exceeds that in Guiwei. Of
31
32 153 19,445 expressed genes, 262 met the registered discovery threshold of genome-wide Benjamini–
33
34 154 Hochberg $q < 0.05$ with an absolute interaction log2 fold change of at least log2(1.5) (Figure
35
36 155 2A); 277 genes met the significance criterion alone and 5,370 met the effect criterion alone. All
37
38
39 156 interaction estimates reported and plotted here are unshrunk maximum-likelihood values;
40
41 157 apeglm shrinkage was used only to define the frozen signature weights. Uniform quality control
42
43 158 then removed 31 candidates for exon-level mappability failure and 25 for gene-model ambiguity,
44
45
46 159 freezing 206 genes as statistical candidates (Figure 2C). Because these quality filters were
47
48 160 applied after significance testing, we verified that the ordering is not load-bearing: restricting the
49
50 161 testing universe to the 13,602 genes that pass uniform quality control and recomputing the
51
52
53 162 Benjamini–Hochberg adjustment over that universe rediscovers all 206 frozen candidates and
54
55 163 adds 12 further genes, so the frozen list is a conservative subset of the QC-first analysis. A post

hoc composite-null sensitivity analysis asked the stricter question whether $|\beta|$ exceeded $\log_2(1.5)$, rather than testing $\beta = 0$ and filtering the estimate afterward. Eight of the 262 statistical candidates passed the DESeq2 composite-null test at $q < 0.05$; the corresponding counts were 7 of 206 QC-retained candidates, 7 of 19 edgeR-gated candidates, and 7 of 16 fully robust candidates. An apeglm false-sign-or-small analysis was less severe, with $s < 0.05$ for 149, 113, 18, and 15 genes in those four nested sets, respectively. These are post hoc sensitivity results, not a redefinition of the registered discovery set. Positive interaction effects (stronger infection response in Yurong1) and negative effects (stronger response in Guiwei) were both well represented. At the pathway level, six Plant Reactome pathways transported through one-to-one *Oryza* orthology met the frozen discovery threshold: activation of the pre-replication complex, assembly of the pre-replication complex, DNA replication initiation, maturation, and circadian rhythm, all with positive enrichment, and jasmonic acid signaling with negative enrichment. Together with the genes, these six pathways, the apeglm-shrunk signed signature weights over the 206 genes, and the conditional transcript-usage results (Section 2.9) were frozen before any external data were opened.

[Insert Figure 2 near here]

2.5 None of the 18 exploratory candidates met the genome-wide interaction criterion

Because the exploratory candidates were selected before any interaction P -values existed, the registered protocol included a formal audit of all 18 against the genome-wide analysis (Figure 2B; Table 2). None met the genome-wide interaction threshold. We then performed the estimand-aligned follow-up requested in review: separate confirmatory DESeq2 infected-versus-mock fits within each cultivar, with Benjamini–Hochberg adjustment over the same 19,445-gene universe in each contrast (Supplementary S14). Thirteen of the 18 legacy genes were significant

in Guiwei and 16 of 18 in Yurong1 at $q < 0.05$; all 13 Guiwei genes and 15 of the 16 Yurong1 genes also exceeded $|\log_2\text{FC}| \geq \log_2(1.5)$. Thus most of the original within-cultivar response claims survive the confirmatory pipeline even though none supports a genome-wide cultivar-by-infection interaction. The strongest interaction-audit candidate, LITCHI001510, reached $q = 0.080$; the EF-hand calcium-binding gene LITCHI019519, whose selected-gene contrast of -3.4 had appeared most dramatic in the exploratory stage, reached only $q = 0.106$ despite a re-estimated interaction effect of -3.33 ; and the median genome-wide interaction q across the 18 candidates was 0.47 . Two candidates, the osmotin/thaumatin-like gene LITCHI009301 and LITCHI028401, additionally failed uniform exon-mappability control, meaning their interaction estimates cannot be attributed unambiguously to a single locus. The audit therefore does not show that these genes are biologically inert. It shows that within-cultivar responsiveness and cultivar-dependent response are distinct claims, and only the latter fails here.

Table 2. Formal genome-wide audit of the 18 exploratory-stage candidates. Interaction effects are re-estimated genome-wide values (Yurong1 response minus Guiwei response, $\log_2$ scale). No candidate met the registered discovery threshold (genome-wide $q < 0.05$ and $|\log_2\text{FC}| \geq \log_2(1.5)$).

Gene	Exploratory annotation	Interaction $\log_2\text{FC}$ (SE)	Genome-wide q	Uniform mappability
LITCHI001510	Poorly characterised responsive gene	-2.83 (0.91)	0.080	Pass
LITCHI019519	EF-hand calcium-binding protein	-3.33 (1.14)	0.106	Pass
LITCHI019299	S-adenosylmethionine-related protein	-1.21 (0.47)	0.175	Pass
LITCHI007102	Phytosulfokine precursor	-0.74 (0.32)	0.238	Pass
LITCHI019183	Plant lipid-transfer protein	$+0.99$ (0.45)	0.273	Pass
LITCHI028401	Selected interaction candidate	-2.41 (1.13)	0.281	Fail
LITCHI012388	Heat-shock protein 70	-0.90 (0.49)	0.381	Pass
LITCHI001552	Cyclase-like stress-related protein	$+0.92$ (0.55)	0.430	Pass
LITCHI001512	Poorly characterised responsive gene	-1.01 (0.63)	0.457	Pass

Gene	Exploratory annotation	Interaction log2FC (SE)	Genome-wide q	Uniform mappability
LITCHI017676	S-adenosylmethionine-related protein	-2.36 (1.53)	0.478	Pass
LITCHI002550	Cytochrome P450	+0.76 (0.96)	0.762	Pass
LITCHI028104	Jacalin-like lectin	+1.39 (1.89)	0.785	Pass
LITCHI002793	Osmotin/thaumatin-like protein	+0.41 (0.55)	0.786	Pass
LITCHI017995	Metalloendoproteinase	+0.48 (0.73)	0.812	Pass
LITCHI009301	Osmotin/thaumatin-like protein	-0.70 (1.08)	0.816	Fail
LITCHI025738	Cyclase-like stress-related protein	-0.22 (0.62)	0.908	Pass
LITCHI011394	Pathogenesis-related protein	-0.25 (0.98)	0.940	Pass
LITCHI018043	Plant lipid-transfer protein	-0.10 (0.52)	0.954	Pass

2.6 Sixteen genes and one pathway passed all internal robustness gates

Each of the 206 frozen genes was required to retain its interaction direction and significance under four predeclared perturbations of the analysis: independent Salmon-based quantification (196 of 206 passed), edgeR quasi-likelihood testing in place of the primary framework (19 passed), a counts-per-million expression filter (205 passed), and all twelve leave-one-library-out refits (125 passed). Eighteen genes passed all four gates jointly, and 16 remained after an additional observed mapping-sensitivity gate that excluded candidates whose effect depended on ambiguously mapped reads (Figure 3A). The alternative-framework gate was by far the most stringent, and the frozen genome-wide edgeR results explain why. The two frameworks agree almost perfectly on estimation: interaction effects correlate at Pearson $r = 0.995$ across all 19,445 genes and $r = 0.999$ across the 206 candidates, sign agreement among the candidates is 100%, and every candidate has nominal edgeR $P < 0.05$. The disagreement is confined to significance calibration. The registered gate required genome-wide edgeR $q < 0.10$ with direction agreement, and edgeR's quasi-likelihood test with robust dispersion estimation admits only 24 genes

1
2
3 218 genome-wide at that threshold (10 at $q < 0.05$), against 262 for the primary Wald analysis at $q <$
4
5 219 0.05. With three libraries per cell, the two dispersion-estimation strategies imply very different
6
7 220 genome-wide significance sets even while agreeing on every point estimate. We therefore
8
9 221 present the results as a hierarchy: 262 statistical candidates, 206 retained after quality control, 19
10
11 222 significant under both frameworks, and 16 passing the complete registered procedure; genes
12
13 223 surviving both frameworks constitute the discovery set's stable core, and the wider 206-gene list
14
15 224 should be read as one framework's candidates rather than as settled discoveries. Of the six frozen
16
17 225 pathways, only circadian rhythm passed its four conjunctive gates (competitive and rotation-
18
19 226 based testing, leading-edge deletion, and an expression- and length-matched null); the other five
20
21 227 all failed the rotation-based self-contained test, and the maturation set additionally failed leading-
22
23 228 edge deletion, indicating enrichment that was not stable across testing frameworks.

24
25
26
27
28
29 229 [Insert Figure 3 near here]

30
31 230 **2.7 The prespecified external evaluation supported two genes and contradicted five**

32
33 231 The primary external estimand, fixed before unlock, was the resistant-minus-susceptible cultivar-
34
35 232 by-infection interaction at 24 h in the PRJNA450886 pericarp time course, tested for each frozen
36
37 233 candidate with Benjamini–Hochberg adjustment across the full candidate-by-contrast family, an
38
39 234 effect threshold of $\log_2(1.5)$, confidence intervals excluding zero, and required agreement with
40
41 235 the discovery direction. Of the 206 frozen genes, 184 were measurable in the external cohort.
42
43 236 Two met every criterion and are cross-context supported: LITCHI005518, annotated as an alpha-
44
45 237 xylosidase involved in xyloglucan remodeling of the cell wall, and LITCHI001963, an
46
47 238 uncharacterized gene with a one-to-one *Oryza* ortholog. Five genes, including a beta-glucosidase
48
49 239 and a Fe(II)/2-oxoglutarate-dependent dioxygenase, passed the magnitude and significance
50
51 240 criteria with the *opposite* direction and are recorded as contradictory; 177 were unsupported at
52
53
54
55
56
57
58
59
60

the prespecified thresholds and 22 not testable (Figure 3C; Table 3). At the level of continuous estimates, discovery and external interaction effects across the 184 measurable candidates were uncorrelated (Pearson $r = -0.05$, 95% CI -0.19 to 0.10 ; Spearman $\rho = -0.04$), with sign concordance of 51.6% (95 of 184), indistinguishable from chance.

The two supported genes are not among the 16 internally robust genes, and none of the 16 internally robust genes achieved external support (Figure 3B). This non-overlap is, by itself, nearly uninformative: for sets of 16 and 2 genes drawn from 206 candidates, the expected overlap under independence is 0.16 genes and the probability of zero overlap is 0.85, so the disjoint sets neither demonstrate nor refute an association between internal stability and cross-context transport. The absence of continuous concordance reported above is the more informative observation. We emphasize that the external cohort differs in tissue (pericarp versus leaf), resistant comparator (Heiye versus Yurong1), and design (time course versus single time point), so absence of support is expected wherever a response is tissue- or genotype-specific, and support that is observed indicates transport across contexts rather than replication of the original experiment. Because Yurong1 and Heiye have not been placed on a common resistance scale relative to Guiwei, we interpret the primary external test as cross-cultivar response transport rather than validation of a resistance phenotype.

Table 3. Cross-context supported and contradictory genes in the primary external evaluation (PRJNA450886, pericarp, 24 h). Discovery effects are Yurong1-minus-Guiwei interaction log2 fold changes in leaf; external effects are Heiye-minus-Guiwei interaction log2 fold changes in pericarp with 95% confidence intervals. All q -values are Benjamini–Hochberg adjusted (discovery: genome-wide; external: across the frozen candidate-by-contrast family).

Gene	Annotation	Discovery log2FC	Discovery q	External log2FC (95% CI)	External q	Outcome
LITCHI005518	Alpha-xylosidase 1	+0.64	0.025	+0.67 (+0.32 to +1.02)	0.007	Supported
LITCHI001963	Uncharacterized; one-to-one <i>Oryza</i>	−1.41	0.024	−0.94 (−1.43 to −0.46)	0.007	Supported

Gene	Annotation	Discovery log2FC	Discovery <i>q</i>	External log2FC (95% CI)	External <i>q</i>	Outcome
	ortholog					
LITCHI030761	Caffeoylshikimate esterase	−0.92	0.040	+1.31 (+0.78 to +1.85)	0.0002	Contradictory
LITCHI029534	Thiamine thiazole synthase 2, chloroplastic	+1.48	0.002	−1.04 (−1.61 to −0.48)	0.010	Contradictory
LITCHI021860	Phosphatidylinositol 4-kinase gamma 2	−0.65	0.032	+0.97 (+0.39 to +1.54)	0.029	Contradictory
LITCHI008313	Beta-glucosidase 44	+0.90	0.045	−1.12 (−1.81 to −0.44)	0.034	Contradictory
LITCHI001700	Fe(II)/2-oxoglutarate-dependent dioxygenase 4	+3.70	0.033	−2.00 (−3.24 to −0.77)	0.035	Contradictory

2.8 Pathway and signature transport was weak and strongly context-dependent

None of the six frozen pathways, including the internally robust circadian-rhythm pathway, passed the external pathway gates in PRJNA450886 (Figure 4A). The frozen signed signature over the 206 discovery genes behaved in a strongly context-dependent way (Figure 4B). At the primary 24-h contrast the signature score was indistinguishable from zero (estimate 0.109, 95% CI −0.079 to 0.297, *q* = 0.349), and at 48 h it was significantly *negative* (−0.240, 95% CI −0.428 to −0.052, *q* = 0.043), a sign reversal relative to discovery. In the generic-transfer study the signature tracked infection positively in leaf (0.107, *q* = 0.0015) but negatively in fruit (−0.181, *q* < 0.001), with a significant leaf-minus-fruit interaction (0.289, *q* < 0.001). The signature therefore carries reproducible infection-associated information, but its sign and magnitude depend on tissue and time in ways that preclude its use as a context-free resistance marker. The exploratory cohort's estimate is reported separately in Figure S3 for completeness and is not confirmatory.

[Insert Figure 4 near here]

2.9 Conditional transcript-usage discovery remained internal

Differential transcript usage was analysed under a conditional gate registered in advance: stage-wise testing had to yield interpretable gene- and transcript-level error control, and no post hoc substitute was permitted if the gate failed. The gate passed, and 225 transcript-usage events across 152 genes were frozen from 15,790 tested transcripts (Figure 5A). External follow-up preserved the role separation of the cohorts (Figure 5B): in the primary cross-context study, 125 events were measurable but unsupported at the prespecified thresholds and 100 were not testable; in the generic-transfer study, annotation incompatibilities left all 225 untestable; and in the exploratory cohort, 4 events were supported and 5 contradictory; we report these outcomes descriptively and do not count them as primary external support, because that cohort's metadata could not be standardised. Transcript-level rewiring between cultivars therefore remains an internally supported, externally unconfirmed layer of the resource.

[Insert Figure 5 near here]

2.10 Orthogonal evidence: informative nulls and reference-ecosystem gaps

The registered orthogonal evidence classes returned uniformly negative or non-testable results, and we report them as such rather than converting them into support. Computational annotation assigned family-level functional labels to 150 of the 206 candidates and left 56 unannotated, but no candidate met the registered high-confidence rule requiring two independent evidence classes with at least 70% supported coverage and no architecture conflict, largely because precomputed InterPro domain architectures were unavailable for all 206 proteins in this non-model species (Figure 6A). Registered promoter-motif discovery tested 927 JASPAR plant position-weight matrices in strand-aware 1-kb and 2-kb windows against 100 expression- and GC-matched genomic background sets. No motif passed in more than 4 of 100 backgrounds against a required

1
2
3 301 80, and none passed FIMO sensitivity (Figure 6B). Small-RNA coherence was not testable
4
5 302 because the frozen PmiREN reference contained no exact *Litchi* entries.
6
7 303 A separate post hoc controlled comparison held the published 18-promoter totals, six explicitly
8
9 304 reported exact-element strings plus one standard PlantCARE TCA-element variant, scoring rule,
10
11 305 and multiplicity correction fixed while changing only the background (Supplementary S15). The
12
13 306 100,000-set randomised-GC rerun closely reproduced the exploratory result. Against 100
14
15 307 expression- and GC-matched genomic promoter sets, ARE lost significance, ABRE gained it,
16
17 308 and MeJA-responsive, TCA, and TC-rich classes remained significant: four of seven classes
18
19 309 passed under either strategy, but the identity of one class changed. Background construction
20
21 310 therefore affected the result but did not by itself explain the exploratory/registered discrepancy.
22
23 311 This comparison is provenance-limited: the archived exploratory package omitted the observed-
24
25 312 motif input TSV, and exact recounting from the current canonical 18 promoters reproduced the
26
27 313 published total only for ARE. Accordingly, it cannot rescue a functional regulatory claim or
28
29 314 establish which foreground representation generated the other published totals. Published studies
30
31 315 reusing any analysed accession, including the primary external study itself [5], remain registered
32
33 316 as prior interpretation rather than independent support.
34
35
36
37
38
39
40

41 317 [Insert Figure 6 near here]

42
43 318 **2.11 Deterministic integration: an empty top tier, twelve Tier B genes, and visible**
44
45 319 **retirements**
46

47
48 320 The registered tier rules integrate all evidence layers without discretion. Tier A required
49
50 321 discovery, complete internal robustness, primary cross-context support, clean mapping and
51
52 322 annotation, and at least one attributable orthogonal class; Tier B required discovery and full
53
54 323 internal robustness with partial or non-testable external evidence; Tier C was reserved for a
55
56
57
58
59
60

robust, externally supported pathway. Applied to the 268 frozen entities (262 statistical gene candidates and 6 pathways), the rules assigned no entity to Tier A or Tier C, 12 genes to Tier B, 191 entities to exploratory status, and 65 to retirement (Figure 6C). The retirements comprise the 56 statistical candidates removed by uniform mapping or gene-model control, the 5 genes contradicted for cross-context directional transport, and 4 candidates with candidate-level observed-mapping or annotation failures. These categories differ in kind, and the released tables preserve the distinction: mapping and annotation failures undermine the discovery itself, whereas an opposite direction in a different tissue and cultivar contrast may reflect genuine context specificity rather than an invalid discovery in the original context. Four internally robust genes, including a probable RPP8-like disease-resistance protein, were measurable but unsupported at the prespecified thresholds in the primary external evaluation and were therefore held at exploratory status rather than promoted. The two cross-context-supported genes likewise remain exploratory because they failed internal robustness. All contradictory and failed outcomes are released alongside the positive ones. We also note that Tier A was difficult to reach by construction in the current reference ecosystem: with the small-RNA class structurally untestable, no motif able to pass the family-wise rule, accession-linked literature excluded as independent support, and the two-class annotation rule unattainable without precomputed domain architectures, the orthogonal-class requirement could not realistically be satisfied by any gene. The empty top tier therefore reflects this structural constraint as much as the expression evidence itself.

The 12 Tier B genes are listed in Table 4. Nine were measurable in the primary external cohort, and all nine were direction-consistent with discovery, although none met the full support criteria. Figure S1 exposes the normalized count of every individual discovery library for these 12 genes,

the 2 cross-context-supported genes, and legacy highlights LITCHI001510 and LITCHI019519; every cultivar-treatment cell contains all three deposited replicates. Their annotations cluster around carbohydrate and cell-wall metabolism (a phosphomannomutase/phosphoglucomutase, a chloroplastic 1,4- α -glucan-branching enzyme, an ERD6-like sugar transporter, a serine carboxypeptidase), membrane transport (a WAT1-related protein), lipid signaling (a DIR1-like lipid-transfer protein, the *Arabidopsis* ortholog of which is required for systemic acquired resistance), programmed-cell-death regulation (a BAG-family chaperone regulator), specialized metabolism (cytochrome P450 81Q32), RNA metabolism (ribonuclease J, exonuclease 1), chromatin-associated regulation (a homeobox-DDT protein), and one entirely unannotated gene.

Table 4. The twelve Tier B genes: internally robust genome-wide statistical candidates with direction-consistent but sub-threshold, or non-measurable, primary external evidence. Discovery effects are interaction log2 fold changes (Yurong1 response minus Guiwei response); positive values indicate a stronger infection response in Yurong1. External columns give the primary 24-h pericarp estimate with 95% confidence interval and family-adjusted *q*; NM denotes genes not measurable in the external cohort.

Gene	Family-level annotation	Discovery log2FC	Genome-wide <i>q</i>	External log2FC (95% CI)	External <i>q</i>
LITCHI014613	Sugar transporter ERD6-like 16	+3.69	8.7×10^{-6}	+0.60 (−0.04 to +1.25)	0.27
LITCHI015982	WAT1-related protein	+3.43	7.8×10^{-7}	+0.77 (+0.07 to +1.48)	0.18
LITCHI005805	Phosphomannomutase/phosphoglucomutase	+2.07	5.5×10^{-8}	NM	—
LITCHI028926	Exonuclease 1	+1.90	5.7×10^{-5}	+0.32 (−0.63 to +1.28)	0.83
LITCHI016077	Serine carboxypeptidase 24	+1.81	4.4×10^{-5}	+0.23 (−1.05 to +1.51)	0.89
LITCHI013915	1,4- α -glucan-branching enzyme 3, chloroplastic	+1.58	4.4×10^{-5}	+0.53 (+0.10 to +0.95)	0.13
LITCHI021546	Homeobox-DDT domain protein RLT3	+1.31	6.6×10^{-5}	+0.12 (−0.30 to +0.54)	0.83
LITCHI022605	Ribonuclease J	+0.74	2.9×10^{-4}	+0.14 (−0.34 to +0.62)	0.83
LITCHI028717	BAG-family molecular chaperone regulator 4	−0.91	1.8×10^{-5}	−1.20 (−2.52 to +0.12)	0.29
LITCHI008721	Unannotated protein	−1.91	7.9×10^{-11}	NM	—

Gene	Family-level annotation	Discovery log2FC	Genome-wide q	External log2FC (95% CI)	External q
LITCHI015740	Putative lipid-transfer protein DIR1	-1.96	1.2×10^{-7}	-0.21 (-1.10 to +0.69)	0.88
LITCHI010877	Cytochrome P450 81Q32	-2.77	2.1×10^{-7}	NM	—

2.12 Parametric simulations quantify the interaction effects detectable with three libraries per cell

The post hoc power analysis simulated negative-binomial counts from the fitted gene-wise means and dispersions while preserving the four-cell discovery design and the 24-h four-cell external design (Figure S2; Supplementary S18). At each of 12 absolute interaction effects from 0 to 4 log2 units, 100 simulations were run sequentially. Genome-wide discovery used 19,445 tests with 262 injected targets per simulation, matching the observed statistical-candidate prevalence; the external analysis adjusted within the 177 frozen genes that passed the simulation's count and dispersion validity filters, with 2 injected targets matching the number of directionally supported genes. Seven additional genes were measurable in the observed 184-gene external contrast but were not fit-eligible for simulation, so the simulated multiplicity family is reported explicitly rather than equated with the observed family. The overall interpolated effect required for 80% detection was 2.44 log2FC (5.4-fold) in discovery and 2.11 log2FC (4.3-fold) in the external candidate-family analysis. At a true |interaction log2FC| of 1.5, overall detection probabilities were 62.7% and 74.0%, respectively. These values are conditional on fitted dispersions treated as known, deposited-library independence, and the simulated non-null prevalence; they quantify the low-power regime but cannot distinguish context dependence from absence of effect for any particular unsupported gene.

1
2
3 404 dramatically in within-cultivar significance while their responses do not differ significantly
4
5 405 between cultivars. Third, two candidates failed uniform-mappability control, a failure mode
6
7 406 invisible to any analysis that trusts a single quantification path. None of this implies the
8
9 407 exploratory candidates are biologically irrelevant; several retain large direction-consistent
10
11 408 estimates and remain reasonable targets for quantitative PCR. Their evidentiary status, however,
12
13 409 is now stated correctly: unconfirmed hypotheses, not validated markers.
14
15
16

17 410 The annotations of the genes that survive scrutiny suggest a recognizable, although still
18
19 411 hypothetical, theme. Annotations related to cell-wall and sugar metabolism appear at every
20
21 412 supported evidence level: among the Tier B genes, a
22
23 413 phosphomannomutase/phosphoglucomutase, a starch-branching enzyme, an ERD6-like sugar
24
25 414 transporter, and a serine carboxypeptidase; and among the two cross-context-supported genes, an
26
27 415 alpha-xylosidase acting on xyloglucan. We emphasize that this is a qualitative clustering of
28
29 416 family-level annotations over small gene sets, assembled from different genes at different
30
31 417 evidence levels, and not a formal over-representation result; no enrichment test was performed,
32
33 418 and several annotations are provisional. The pattern is nonetheless consistent with the emphasis
34
35 419 on sugar metabolism and cell-wall reinforcement in prior comparisons of resistant and
36
37 420 susceptible lychee cultivars [5], and with the general role of apoplastic sugar status and wall
38
39 421 remodeling in oomycete resistance. The Tier B set also contains a DIR1-like lipid-transfer
40
41 422 protein, whose *Arabidopsis* counterpart is a mobile signal in systemic acquired resistance, and a
42
43 423 BAG-family chaperone regulator; BAG proteins act as HSP70 co-chaperones within the
44
45 424 chaperone networks that shape plant immune responses and immune-associated cell death [9].
46
47 425 The strong negative discovery enrichment of jasmonic acid signaling, although it did not survive
48
49 426 external evaluation, is likewise a plausible axis of cultivar difference given the jasmonate-
50
51
52
53
54
55
56
57
58
59
60

1
2
3 427 associated promoter elements suggested in the exploratory stage and the known antagonism
4
5 428 between jasmonate and salicylate defence programs. We advance all of these as prioritised,
6
7 429 testable hypotheses; the tier system exists precisely to prevent us from advancing them as more.
8
9
10 430 The external evaluation delivers a second structural lesson. The internally robust and externally
11
12 431 supported sets did not overlap, but this disjunction carries almost no information on its own: with
13
14 432 16 and 2 genes selected from 206 candidates, the expected overlap under independence is 0.16
15
16 433 and zero overlap occurs with probability 0.85. The informative observations are continuous ones:
17
18 434 discovery and external interaction effects were uncorrelated across the 184 measurable
19
20 435 candidates, the frozen signature reversed sign between 24 h and 48 h in the external time course,
21
22 436 and the same signature tracked infection in opposite directions in leaf and fruit of the transfer
23
24 437 cohort. The straightforward reading is that the cultivar-dependent component of the early
25
26 438 infection response is substantially tissue- and time-specific, so that a candidate's failure to
27
28 439 transport from leaf at 24 h to pericarp at 48 h is expected behaviour rather than evidence against
29
30 440 the discovery. This has a practical corollary for breeding applications: expression-based
31
32 441 resistance markers derived from one tissue and time point should not be assumed portable, and
33
34 442 marker development for this pathosystem will require validation in the tissue and developmental
35
36 443 window where deployment is intended.
37
38 444 The orthogonal layers contribute mainly disciplined null results, and we consider their inclusion
39
40 445 one of the more useful aspects of the resource. The post hoc controlled PlantCARE comparison
41
42 446 now isolates the background rule as far as the archived inputs permit. It shows that genomic
43
44 447 expression/GC matching changes which class passes (ARE is replaced by ABRE) but does not
45
46 448 eliminate the MeJA, TCA, or TC-rich signals when the published foreground totals are held
47
48 449 fixed. Background choice is therefore not a sufficient explanation for the failure of the registered
49
50
51
52
53
54
55
56
57
58
59
60

JASPAR analysis; the changed foreground, motif representation, scoring, and robustness rule remain consequential. Because the original observed-motif TSV was not archived and most published totals cannot be reproduced from the current canonical promoters, the controlled result is a provenance audit rather than evidence of regulatory function. The annotation and small-RNA outcomes likewise expose reference-ecosystem gaps (no precomputed InterPro architectures and no *Litchi* entries in the frozen miRNA reference) that bound what any computational study of this species can currently claim.

Several limitations qualify all of the above. The deposited designs are small, with three libraries per cell, and the independence of those libraries at the level of source trees, pooling, and extraction is unreported, so our q -values are conditional on deposited-library independence. The public metadata also leave the inoculation protocol underspecified, including the inoculum form (mycelia, sporangia, or zoospores), the developmental stage of the sampled leaves, and whether the tissue was wounded before inoculation, so all conclusions are additionally conditional on the depositors' experimental design. All quantification depends on a single reference assembly, and although uniform and observed mapping checks removed the clearest failure modes, residual reference bias cannot be excluded. The primary external cohort is deliberately heterogeneous with respect to the discovery cohort, which makes support conservative and non-support ambiguous; no genuinely independent replication cohort (same tissue, same time, different biological material) exists in the public record for this design. Unsupported external outcomes additionally conflate limited power, context dependence, and true absence of effect. The simulations quantify the first component: conditional 80% minimum detectable effects were 2.44 log₂FC (5.4-fold) for genome-wide discovery and 2.11 log₂FC (4.3-fold) for candidate-family external evaluation, with still poorer performance among low-expression genes. These

1

2

3

473 conditional curves do not identify which unsupported genes are context-specific or truly null.

5

474 And the entire study is computational: no candidate, including the Tier B set, has been perturbed

6

7

475 experimentally. The natural next steps are quantitative PCR of the Tier B genes in independently

8

9

476 infected material of both discovery cultivars, functional tests of the alpha-xylosidase and DIR1-

10

11

477 like candidates, and, for any regulatory hypothesis, direct assays such as promoter-reporter or

12

13

478 chromatin-binding experiments rather than further motif counting.

14

15

16

17

18

479 **4. Conclusions**

19

20

21

480 By carrying a single pathosystem from exploratory candidate discovery through a prospectively

22

23

481 registered genome-wide discovery framework with prespecified cross-context external

24

25

482 evaluation, this study establishes which claims about cultivar-dependent responses of lychee to

26

27

483 *P. litchii* the current public data can and cannot support. It can support a set of 206 QC-retained

28

29

484 genome-wide statistical interaction candidates, of which 19 are significant under both statistical

30

31

485 frameworks, a stable core of 16 internally robust genes, 12 Tier B candidates whose annotations

32

33

486 include cell-wall, sugar-metabolism, lipid-signaling, and RNA-metabolism functions, and 2

34

35

487 genes with cross-context support in an independent tissue and cultivar contrast. It cannot support

36

37

488 the original 18 exploratory candidates as validated markers, any pathway- or signature-level

38

39

489 claim of context-free transport, or any regulatory-element claim under matched-background

40

41

490 testing. All frozen results, including nulls, contradictions, and retirements, are released with the

42

43

491 analysis code and checksummed source data, so that the next stage of experimental validation

44

45

492 can begin from an honest baseline.

46

47

48

49

50

51

52

53

54

55

56

57

58

59

60

5. Materials and Methods

5.1 Study design, registration, and evidence vocabulary

The confirmatory stage was governed by a time-stamped protocol, with amendments logged, distributed with the accompanying repository. The registration chronology is as follows. The exploratory stage was completed first, so the discovery dataset had already been analysed, under different models, before registration. Protocol version 1.0.0 was frozen on 18 August 2026, with its SHA-256 hash (ef668486...) recorded alongside the protocol text, before any confirmatory expression outcome was computed. The amendment log records 20 outcome-blind amendments (clarifications, implementation corrections, and optimisations, all time-stamped before any discovery or external expression outcome was viewed) and 3 post-outcome amendments confined to technical quality control and manifest verification, each marked as such. The primary external study had been published previously [5], and its headline biological conclusions were known to the analyst before registration; dataset roles and all thresholds were nevertheless fixed before any external count matrix was processed. Because all cohorts are public, the staged external unlock is a workflow-enforced computational safeguard, not blinding in the experimental sense. The protocol fixed dataset roles (Table 1), statistical models, discovery and evaluation thresholds, conditional gates, empty-set behaviour, and the vocabulary used throughout this paper: *frozen* denotes a computed result fixed at a recorded checkpoint and never modified afterward; *discovery* refers only to the registered genome-wide interaction analysis in PRJNA830488; *robustness* refers only to predeclared sensitivity analyses within that cohort; *cross-context support* denotes a passing frozen external test in a cohort that differs in tissue, comparator, or design, and is never described as replication; and *orthogonal* evidence classes (annotation, motifs, small RNA, literature) are reported separately from expression support and

1
2
3 516 cannot substitute for it. No external outcome could modify frozen discovery membership,
4
5 517 signature weights, or thresholds.
6
7

8
9 518 **5.2 Exploratory-stage analyses (Stage 1)**
10

11 519 The exploratory stage analysed the 12 GSE201243 libraries with FastQC v0.11.9 quality control,
12
13 520 Trimmomatic v0.39 trimming, TopHat v2.0.12/Bowtie2 v2.2.3 alignment to the
14
15 521 SCAU_Lch_v2.0 assembly (GCA_019925255.1), HTSeq v2.0.2 union-mode counting, and
16
17 522 DESeq2 v1.42.0 per-cultivar Wald contrasts (design ~ treatment) with Benjamini–Hochberg
18
19 523 adjustment at $P < 0.05$ [10–15]. Candidate prioritisation combined significance, fold-change
20
21 524 magnitude, and annotation plausibility; structural, conserved-motif, and promoter analyses used
22
23 525 TBtools-II v2.096, MEME v5.5.5, and PlantCARE [16–18]; and the promoter-background
24
25 526 comparison tested observed *cis*-element counts in the 18 candidate promoters against 100,000
26
27 527 randomised sets of eighteen 2-kb sequences with lychee-like GC content, scoring exact motifs
28
29 528 and reverse complements, with one-sided empirical P -values adjusted by Benjamini–Hochberg
30
31 529 [19]. The selected-gene interaction contrast was computed as $\log_2\text{FC}(\text{Yurong1, infected versus}$
32
33 530 $\text{mock}) - \log_2\text{FC}(\text{Guiwei, infected versus mock})$ over the prioritised genes and was explicitly
34
35 531 labeled a prioritisation effect size rather than a genome-wide test.
36
37
38
39
40
41

42 532 **5.3 Data provenance and eligibility**
43

44 533 National Centre for Biotechnology Information and Gene Expression Omnibus metadata for all
45
46 534 five cohorts were reconciled into a biological-unit registry (Supplementary S1). The original
47
48 535 repository preserved accession decisions but not the historical query log, so search coverage was
49
50 536 reconstructed retrospectively on 27 August 2026 and is not represented as prospective
51
52
53 537 provenance. A broad NCBI GEO series query returned six records and a corresponding
54
55
56
57
58
59
60

BioProject query returned seven; a narrow SRA cross-check returned 42 experiment-level records. After collapsing the GSE222652 super-series, excluding PRJNA1157370 because it used *Colletotrichum gloeosporioides*, and excluding PRJNA268587/GSE63658 because it studied fruit senescence/storage without *P. litchii* inoculation, the five prespecified cohorts remained (Supplementary S16a–b). Because source-tree, pooling, harvest, and extraction independence were not reported for the discovery libraries, all inferential q -values are conditional on deposited-library independence. The primary model required at least three included libraries in every cultivar-treatment cell.

5.4 Confirmatory preprocessing and quality gates

Host (SCAU_Lch_v2.0) and pathogen references were combined with prefixed contig names to absorb pathogen-derived reads. Paired reads were checked with FastQC, trimmed with fastp [20], aligned with STAR two-pass mapping [21], counted at the gene level with featureCounts [22], and independently quantified with Salmon [23]. Libraries were excluded if fewer than 10 million read pairs survived trimming or fewer than 40% of reads aligned uniquely; all 12 discovery libraries passed. GenMap-based exon-uniqueness scores [24] and candidate-level observed-mapping checks were fixed before discovery and used for the uniform mapping and observed mapping-sensitivity gates.

5.5 Genome-wide discovery model

DESeq2 [15] fitted gene counts with the design $\sim$ cultivar + treatment + cultivar:treatment (reference levels Guiwei, mock). The primary coefficient, cultivarYurong1:treatmentinfected, estimates the difference in infected-minus-mock response between Yurong1 and Guiwei. Genes required counts of at least 10 in at least three libraries; statistical candidates required genome-

1
2
3 560 wide Benjamini–Hochberg $q < 0.05$ and $|\log_2\text{FC}| \geq \log_2(1.5)$ [19]. The registered rule tests the
4
5 561 point null of zero interaction and applies the effect threshold to the observed estimate. Post hoc
6
7 562 sensitivity used DESeq2 with `lfcThreshold = log2(1.5)` and `altHypothesis =`
8
9 `greaterAbs`, plus `apegglm` false-sign-or-small s -values for the same threshold; neither analysis
10 563
11 modified frozen membership. To align the legacy audit with its original estimand, separate post
12 564
13 hoc DESeq2 models (`~ treatment`) were fitted within Guiwei and Yurong1, using the same
14 565
15 count matrix, 19,445-gene universe, and genome-wide Benjamini–Hochberg adjustment within
16 566
17 each cultivar. Unshrunk maximum-likelihood estimates are reported throughout; `apegglm`-
18 567
19 shrunk interaction estimates [25] were used only to define the signed signature weights over
20 568
21 the 206 frozen genes. Mappability and gene-model metrics were computed for all 19,445 tested
22 569
23 genes; as a sensitivity check, recomputing the Benjamini–Hochberg adjustment over only the
24 570
25 13,602 QC-eligible genes rediscovered all 206 frozen candidates plus 12 additional genes,
26 571
27 confirming that applying quality control after significance testing was conservative here.
28 572
29
30
31
32
33

34 573 **5.6 Pathway and transcript-usage discovery**

35
36 574 Plant Reactome pathway memberships [26] were transported to lychee through one-to-one
37 575
38 reciprocal-best DIAMOND protein matches to *Oryza* [27], and pathway discovery applied the
39 576
40 frozen genome-wide interaction statistic to this fixed mapping with camera, roast, and fgsea
41 577
42 testing plus leading-edge-deletion and expression/length-matched null gates for robustness [28–
43 578
44 30]. Differential transcript usage used Salmon transcript abundances analysed with DRIMSeq
45 579
46 and DEXSeq under stageR stage-wise error control [31–33]; the conditional inferential gate for
47 580
48 interpretable two-level error control was registered in advance, and no substitute analysis was
49 581
50 permitted had it failed.
51
52
53
54
55
56
57
58
59
60

5.7 Internal robustness gates

Each frozen gene was retested under Salmon-based gene-level quantification, edgeR quasi-likelihood interaction testing [34], a counts-per-million expression filter, and all leave-one-library-out refits, with predeclared criteria combined conjunctively; no additive scoring was used. The Salmon gate required direction agreement and an absolute effect difference of at most 0.5 log₂ units; the edgeR analysis used the identical design matrix and interaction coefficient with TMM normalization, robust dispersion estimation (estimateDisp and glmQLFit with robust settings), and glmQLFTest, with Benjamini–Hochberg adjustment computed genome-wide over all 19,445 genes and a gate of genome-wide $q < 0.10$ plus direction agreement; the expression-filter gate refit the primary model after a counts-per-million filter; and the leave-one-library-out gate required direction agreement in all 12 refits and $q < 0.10$ in at least 10 of 12. The observed mapping-sensitivity gate removed candidates whose effect depended on ambiguously mapped reads. Frozen pathways required agreement across camera, roast, leading-edge deletion, and the matched-null comparison.

5.8 External frozen evaluation

PRJNA450886 was analysed with the design $\sim \text{cultivar} \times \text{treatment} \times \text{time}$, with time encoded as a categorical factor and 24 h as the reference level, so that the cultivar-by-treatment coefficient estimates the 24-h interaction and the three-way terms give the 6-h and 48-h contrasts; the primary estimand was the resistant-minus-susceptible cultivar-by-infection interaction at 24 h, with 6-h and 48-h contrasts secondary. Frozen candidate tests used Benjamini–Hochberg adjustment across the full candidate-by-contrast family, an absolute effect threshold of log₂(1.5), confidence intervals excluding zero, and discovery-direction agreement; opposite-direction passes were recorded as contradictory. PRJNA922966 was analysed with $\sim \text{tissue} + \text{treatment} +$

1
2
3 605 tissue:treatment for generic-transfer tests, and PRJNA1090613 outcomes were confined to
4
5 606 exploratory status by protocol. The frozen signed signature was evaluated as a library-level score
6
7
8 607 with the same family-wise adjustment.
9

10
11 608 **5.9 Annotation, motif, small-RNA, and literature classes**
12

13 609 Representative proteins were chosen by longest coding sequence with lexicographic tie-breaking.
14
15 610 Reviewed Swiss-Prot (release 2026_02) DIAMOND matches [35], targeted InterPro/Pfam
16
17 611 domain assignments [36], and one-to-one *Oryza* orthology formed three separate annotation
18
19 612 classes; the high-confidence label required at least two classes, at least 70% supported coverage,
20
21 613 and no architecture conflict. Promoter-motif analysis tested 927 JASPAR plant position-weight
22
23 614 matrices [37] in strand-aware 1-kb and 2-kb promoter windows against 100 expression- and GC-
24
25 615 matched background sets, with AME as the primary test, FIMO as sensitivity [38,39], robustness
26
27 616 requiring passes in at least 80 of 100 backgrounds in both windows, and cognate transcription-
28
29 617 factor expression required for attribution. The PmiREN bulk archive [40] was audited before
30
31 618 small-RNA analysis; because it contained no exact *Litchi* entries, the class was frozen as not
32
33 619 testable and no study-derived reference was substituted. Published studies reusing any analysed
34
35 620 accession were registered as prior interpretation. For the post hoc controlled PlantCARE
36
37 621 comparison, published observed totals across the 18 legacy promoters were fixed and tested
38
39 622 against (i) 100,000 simulated sets of eighteen 2-kb sequences at 34% GC and (ii) 100 genomic
40
41 623 sets of eighteen promoters matched by mean-expression and promoter-GC quintiles with the
42
43 624 same seeded minimum-distance rule used by the registered motif pipeline. The six reported
44
45 625 motifs and one standard TCA variant, together with reverse complements, were counted and one-
46
47 626 sided empirical *P*-values were adjusted across seven classes. The missing archived observed-
48
49
50
51
52
53
54
55
56
57
58
59
60

motif TSV and canonical-promoter recount discrepancies were retained as explicit provenance limitations (Supplementary S15).

5.10 Simulation-based power analysis

Negative-binomial parametric simulations used DESeq2 maximum-a-posteriori gene dispersions and fitted sample means from the discovery counts and from the 24-h subset of PRJNA450886. Fitted cultivar and treatment main effects and size factors were retained, background-gene interaction coefficients were set to zero, and equal numbers of positive and negative effects were injected at 12 absolute log2FC values (0, 0.5, 0.75, 1, 1.25, 1.5, 1.75, 2, 2.5, 3, 3.5, 4). Each value had 100 simulations. Discovery injected 262 targets and applied Benjamini–Hochberg correction over 19,445 genes plus the registered effect filter. External simulation injected 2 targets among the 177 frozen genes passing the simulation count and dispersion validity filters and additionally required the prespecified sign and a 95% interval excluding zero; the observed external table contained 184 measurable genes, seven of which were not fit-eligible for the simulation. Dispersions were treated as known during simulated Wald refits, making the curves conditional parametric power rather than a claim about population-level biological replication. Eighty-percent minimum detectable effects were linearly interpolated between grid points, with expression-quartile curves and Wilson intervals reported in Figure S2.

5.11 Tier rules, software versions, and reproducibility

Tier A required discovery, complete internal robustness, primary cross-context support, clean mapping and annotation, and at least one attributable orthogonal class; Tier B required discovery and complete internal robustness with partial or non-testable primary external evidence; Tier C was reserved for a robust, externally supported pathway. Contradictory external direction or

1
2
3 649 mapping/annotation failure retired an entity, and an empty Tier A was an allowed outcome. All
4
5 650 stages ran under Snakemake 9.25.2 workflows [41] enforcing staged discovery, a controlled
6
7 651 external unlock, verified cleanup, and SHA-256 manifests over inputs and outputs. Confirmatory
8
9 652 executables were FastQC 0.12.1, fastp 1.3.6, STAR 2.7.10b, featureCounts/Subread 2.1.1,
10
11 653 Salmon 2.5.1, MEME Suite/AME/FIMO 5.5.9, and DIAMOND 2.2.5; GenMap 1.3.0 is
12
13 654 recoverable from the pinned specification although its executable is absent from the
14
15 655 reconstructed environment. Statistical work used R 4.3.3/Bioconductor 3.18 with DESeq2
16
17 656 1.42.0, apegglm 1.24.0, edgeR 4.0.16, DRIMSeq 1.30.0, DEXSeq 1.48.0, stageR 1.24.0, and
18
19 657 BiocParallel 1.36.0 (Supplementary S17 records the evidence status of every version). The
20
21 658 resolved environment, synthetic fixtures, per-figure source data, and amendment log are
22
23 659 distributed with the repository. All six main figures are regenerated from frozen result tables by a
24
25 660 script that recomputes and asserts every headline count.
26
27
28
29
30
31

32 661 **Ethics**

33
34

35 662 This study reanalysed publicly available sequencing datasets and involved no new experiments
36
37 663 with humans, human tissue, animals or field sampling; ethical approval was not required.
38
39
40

41 664 **Acknowledgements**

42
43

44 665 The author has no acknowledgements to declare.
45
46

47 666 **Funding**

48
49

50 667 This work received no external funding.
51
52
53
54
55
56
57
58
59
60

668 **Authors' contributions**

669 Eric Zhuang: conceptualization, data curation, formal analysis, investigation, methodology,
670 project administration, software, validation, visualization, writing--original draft and writing--
671 review and editing. The author approved the submitted version and is accountable for the work.

672 **Competing interests**

673 The author declares no conflicts of interest.

674 **Data accessibility**

675 All analysed sequencing data are public under PRJNA830488/GSE201243, PRJNA450886,
676 PRJNA922966/GSE222651, PRJNA922965/GSE222650, and PRJNA1090613/GSE262200.
677 Supplementary Tables S1-S18, Figures S1-S3, and tab-separated figure source data are archived
678 under CC BY 4.0 at <https://doi.org/10.5281/zenodo.22240717> [42]. The analysis source code,
679 Snakemake workflows, configuration files, tests and environment specifications are supplied as
680 electronic supplementary material in RSOS_analysis_code.zip. Together, these public records
681 and electronic supplementary files provide the data, code and supporting materials required to
682 reproduce the reported results.

683 **Accession numbers**

684 Sequence data analysed in this article can be found in the NCBI SRA and GEO repositories
685 under PRJNA830488 (GSE201243), PRJNA450886, PRJNA922966 (GSE222651),
686 PRJNA922965 (GSE222650), and PRJNA1090613 (GSE262200). Supplementary data are
687 permanently archived at <https://doi.org/10.5281/zenodo.22240717>.

1

2

3

4

5

6

7

8

9

10

11

12

13

14

15

16

17

18

19

20

21

22

23

24

25

26

27

28

29

30

31

32

33

34

35

36

37

38

39

40

41

42

43

44

45

46

47

48

49

50

51

52

53

54

55

56

57

58

59

60

688

689

690

691

692

693

694

695

696

697

698

699

700

701

702

703

704

705

706

707

708

709

710

711

712

713

714

715

716

717

718

719

Use of artificial intelligence

Generative AI (OpenAI Codex) was used to assist with language editing and journal-specific document formatting. It was not used to generate data, perform analyses, interpret results or draw scientific conclusions. The author reviewed and takes responsibility for all content.

Supplementary material

The following items are archived in the electronic supplementary material at <https://doi.org/10.5281/zenodo.22240717>:

- Table S1. Biological-unit registry.
- Table S2. Per-library quality control.
- Table S3. Genome-wide discovery statistics.
- Table S4. Robustness file manifest.
- Table S5. All external frozen tests.
- Table S6. Pathway and signature tests.
- Table S7. Conditional differential-transcript-usage results.
- Table S8. Candidate annotation and orthology evidence.
- Table S9. Small-RNA reference-gate outcome.
- Table S10. Motif-background tests and sensitivity results.
- Table S11. Accession-aware evidence registry.
- Table S12. Scripts, environments, and command inventory.
- Table S13. Protocol amendment and deviation log.
- Table S14. Legacy within-cultivar audit.
- Table S15. Controlled promoter-motif background comparison.
- Table S16a. Reconstructed dataset-search queries.
- Table S16b. Dataset eligibility decisions.
- Table S17. Exact software versions.
- Table S18. Power-simulation minimum detectable effects.
- Figure S1. Replicate-level normalized counts.
- Figure S2. Conditional parametric power curves.
- Figure S3. Exploratory signed-signature estimate.
- Data S1. Tab-separated source data for all main and supplementary analytical figures.
- Code S1. Analysis code, workflows, configurations, tests, metadata, and environment specifications.

References

1. Morton JF. Lychee. In: Fruits of Warm Climates. Winterville: Creative Resource Systems; 1987. p. 249–259.
2. Sun J, Gao Z, Zhang X, Zou X, Cao L, Wang J. Transcriptome analysis of *Phytophthora litchii* reveals pathogenicity arsenals and confirms taxonomic status. PLoS ONE. 2017;12:e0178245. doi:10.1371/journal.pone.0178245.
3. Yi C, Jiang Y, Shi J, Qu H, Xue S, Duan X, et al. ATP-regulation of antioxidant properties and phenolics in litchi fruit during browning and pathogen infection process. Food Chemistry. 2010;118:42–47. doi:10.1016/j.foodchem.2009.04.074.
4. Hu G, Feng J, Xiang X, Wang J, Salojärvi J, Liu C, et al. Two divergent haplotypes from a highly heterozygous lychee genome suggest independent domestication events for early and late-maturing cultivars. Nature Genetics. 2022;54:73–83. doi:10.1038/s41588-021-00971-3.
5. Sun J, Cao L, Li H, Wang G, Wang S, Li F, et al. Early responses given distinct tactics to infection of *Peronophythora litchii* in susceptible and resistant litchi cultivar. Scientific Reports. 2019;9:2810. doi:10.1038/s41598-019-39100-w.
6. Liu H, Yan Q, Jiang Y, Shi F, Chen J, Cai C, et al. Identification of LcCDPKs and analysis of their expression patterns in response to downy mildew stresses in lychee. Journal of Fruit Science. 2023;40:442–456. doi:10.13925/j.cnki.gsxb.20220307.
7. Li P, Li W, Zhou X, Situ J, Xie L, Xi P, et al. *Peronophythora litchii* RXLR effector PlAvh202 destabilizes a host ethylene biosynthesis enzyme. Plant Physiology. 2023;193:756–774. doi:10.1093/plphys/kiad311.
8. National Center for Biotechnology Information. GEO accession GSE201243. <https://www.ncbi.nlm.nih.gov/geo/query/acc.cgi?acc=GSE201243>.
9. Berka M, Kopecká R, Berková V, Brzobohatý B, Černý M. Regulation of heat shock proteins 70 and their role in plant immunity. Journal of Experimental Botany. 2022;73:1894–1909. doi:10.1093/jxb/erab549.
10. Andrews S. FastQC: a quality control tool for high throughput sequence data. Babraham Bioinformatics; 2010.
11. Bolger AM, Lohse M, Usadel B. Trimmomatic: a flexible trimmer for Illumina sequence data. Bioinformatics. 2014;30:2114–2120. doi:10.1093/bioinformatics/btu170.
12. Kim D, Pertea G, Trapnell C, Pimentel H, Kelley R, Salzberg SL. TopHat2: accurate alignment of transcriptomes in the presence of insertions, deletions and gene fusions. Genome Biology. 2013;14:R36. doi:10.1186/gb-2013-14-4-r36.
13. Anders S, Pyl PT, Huber W. HTSeq—a Python framework to work with high-throughput sequencing data. Bioinformatics. 2015;31:166–169. doi:10.1093/bioinformatics/btu638.

1
2
3 755 14. Genome Warehouse. *Litchi chinensis* SCAU_Lch_v2.0 assembly (GCA_019925255.1).
4 756 National Genomics Data Center, China National Center for Bioinformation.
5
6 757 15. Love MI, Huber W, Anders S. Moderated estimation of fold change and dispersion for RNA-
7 758 seq data with DESeq2. *Genome Biology*. 2014;15:550. doi:10.1186/s13059-014-0550-8.
8
9 759 16. Chen C, Chen H, Zhang Y, Thomas HR, Frank MH, He Y, et al. TBtools: an integrative
10 760 toolkit developed for interactive analyses of big biological data. *Molecular Plant*.
11 761 2020;13:1194–1202. doi:10.1016/j.molp.2020.06.009.
12
13 762 17. Bailey TL, Elkan C. Fitting a mixture model by expectation maximization to discover motifs
14 763 in biopolymers. In: *Proceedings of the Second International Conference on Intelligent*
15 764 *Systems for Molecular Biology*. Menlo Park: AAAI Press; 1994. p. 28–36.
16
17 765 18. Lescot M, Déhais P, Thijs G, Marchal K, Moreau Y, Van de Peer Y, et al. PlantCARE, a
18 766 database of plant cis-acting regulatory elements and a portal to tools for in silico analysis of
19 767 promoter sequences. *Nucleic Acids Research*. 2002;30:325–327. doi:10.1093/nar/30.1.325.
20
21 768 19. Benjamini Y, Hochberg Y. Controlling the false discovery rate: a practical and powerful
22 769 approach to multiple testing. *Journal of the Royal Statistical Society, Series B*. 1995;57:289–
23 770 300.
24
25 771 20. Chen S, Zhou Y, Chen Y, Gu J. fastp: an ultra-fast all-in-one FASTQ preprocessor.
26 772 *Bioinformatics*. 2018;34:i884–i890. doi:10.1093/bioinformatics/bty560.
27
28 773 21. Dobin A, Davis CA, Schlesinger F, Drenkow J, Zaleski C, Jha S, et al. STAR: ultrafast
29 774 universal RNA-seq aligner. *Bioinformatics*. 2013;29:15–21.
30 775 doi:10.1093/bioinformatics/bts635.
31
32 776 22. Liao Y, Smyth GK, Shi W. featureCounts: an efficient general purpose program for assigning
33 777 sequence reads to genomic features. *Bioinformatics*. 2014;30:923–930.
34 778 doi:10.1093/bioinformatics/btt656.
35
36 779 23. Patro R, Duggal G, Love MI, Irizarry RA, Kingsford C. Salmon provides fast and bias-aware
37 780 quantification of transcript expression. *Nature Methods*. 2017;14:417–419.
38 781 doi:10.1038/nmeth.4197.
39
40 782 24. Pockrandt C, Alzamel M, Iliopoulos CS, Reinert K. GenMap: ultra-fast computation of
41 783 genome mappability. *Bioinformatics*. 2020;36:3687–3692.
42 784 doi:10.1093/bioinformatics/btaa222.
43
44 785 25. Zhu A, Ibrahim JG, Love MI. Heavy-tailed prior distributions for sequence count data:
45 786 removing the noise and preserving large differences. *Bioinformatics*. 2019;35:2084–2092.
46 787 doi:10.1093/bioinformatics/bty895.
47
48 788 26. Naithani S, Gupta P, Preece J, D'Eustachio P, Elser J, Garg P, et al. Plant Reactome
49 789 knowledgebase: empowering plant pathway exploration and OMICS data analysis. *Nucleic*
50 790 *Acids Research*. 2024;52:D1538–D1547. doi:10.1093/nar/gkad1052.
51
52
53
54
55
56
57
58
59
60

- 791 27. Buchfink B, Xie C, Huson DH. Fast and sensitive protein alignment using DIAMOND.
792 Nature Methods. 2015;12:59–60. doi:10.1038/nmeth.3176.
- 793 28. Wu D, Smyth GK. Camera: a competitive gene set test accounting for inter-gene correlation.
794 Nucleic Acids Research. 2012;40:e133. doi:10.1093/nar/gks461.
- 795 29. Wu D, Lim E, Vaillant F, Asselin-Labat ML, Visvader JE, Smyth GK. ROAST: rotation
796 gene set tests for complex microarray experiments. Bioinformatics. 2010;26:2176–2182.
797 doi:10.1093/bioinformatics/btq401.
- 798 30. Korotkevich G, Sukhov V, Budin N, Shpak B, Artyomov MN, Sergushichev A. Fast gene set
799 enrichment analysis. bioRxiv. 2021:060012. doi:10.1101/060012.
- 800 31. Nowicka M, Robinson MD. DRIMSeq: a Dirichlet-multinomial framework for multivariate
801 count outcomes in genomics. F1000Research. 2016;5:1356.
802 doi:10.12688/f1000research.8900.2.
- 803 32. Anders S, Reyes A, Huber W. Detecting differential usage of exons from RNA-seq data.
804 Genome Research. 2012;22:2008–2017. doi:10.1101/gr.133744.111.
- 805 33. Van den Berge K, Sonesson C, Robinson MD, Clement L. stageR: a general stage-wise
806 method for controlling the gene-level false discovery rate in differential expression and
807 differential transcript usage. Genome Biology. 2017;18:151. doi:10.1186/s13059-017-1277-
808 0.
- 809 34. Robinson MD, McCarthy DJ, Smyth GK. edgeR: a Bioconductor package for differential
810 expression analysis of digital gene expression data. Bioinformatics. 2010;26:139–140.
811 doi:10.1093/bioinformatics/btp616.
- 812 35. The UniProt Consortium. UniProt: the Universal Protein Knowledgebase in 2023. Nucleic
813 Acids Research. 2023;51:D523–D531. doi:10.1093/nar/gkac1052.
- 814 36. Jones P, Binns D, Chang HY, Fraser M, Li W, McAnulla C, et al. InterProScan 5: genome-
815 scale protein function classification. Bioinformatics. 2014;30:1236–1240.
816 doi:10.1093/bioinformatics/btu031.
- 817 37. Rauluseviciute I, Riudavets-Puig R, Blanc-Mathieu R, Castro-Mondragon JA, Ferenc K,
818 Kumar V, et al. JASPAR 2024: 20th anniversary of the open-access database of transcription
819 factor binding profiles. Nucleic Acids Research. 2024;52:D174–D182.
820 doi:10.1093/nar/gkad1059.
- 821 38. McLeay RC, Bailey TL. Motif enrichment analysis: a unified framework and an evaluation
822 on ChIP data. BMC Bioinformatics. 2010;11:165. doi:10.1186/1471-2105-11-165.
- 823 39. Grant CE, Bailey TL, Noble WS. FIMO: scanning for occurrences of a given motif.
824 Bioinformatics. 2011;27:1017–1018. doi:10.1093/bioinformatics/btr064.
- 825 40. Guo Z, Kuang Z, Wang Y, Zhao Y, Tao Y, Cheng C, et al. PmiREN: a comprehensive
826 encyclopedia of plant miRNAs. Nucleic Acids Research. 2020;48:D1114–D1121.
827 doi:10.1093/nar/gkz894.

1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60

41. Mölder F, Jablonski KP, Letcher B, Hall MB, Tomkins-Tinch CH, Sochat V, et al.
Sustainable data analysis with Snakemake. F1000Research. 2021;10:33.
doi:10.12688/f1000research.29032.2.

42. Zhuang E. 2026 Supplementary material for Cultivar-dependent transcriptional responses of
lychee to *Peronophythora litchii*: a registered genome-wide analysis. Zenodo.
doi:10.5281/zenodo.22240717.

Figure Legends

Figure 1. Study cohorts, discovery quality control, and expression structure. (A) The five public cohorts with prespecified roles and deposited library counts. (B) Per-library technical quality in the discovery cohort: fastp read survival, STAR unique-alignment rate, and Salmon mapping rate; all 12 libraries passed the fixed inclusion gates. (C) Principal-component analysis of the discovery libraries separates cultivar and infection status.

Figure 2. Genome-wide discovery and formal reassessment of the exploratory candidates. (A) Cultivar-by-infection interaction effects across 19,445 expressed genes; 262 genes met the statistical threshold and 206 were retained after uniform mapping and gene-model quality control. (B) Genome-wide re-estimation of the 18 exploratory (legacy) candidates with 95% confidence intervals; orange marks uniform-mappability failures. No legacy candidate met the genome-wide interaction criterion. (C) Evidence attrition from tested genes to internally robust candidates.


Figure 3. Internal robustness and the primary external evaluation. (A) Candidates passing each predeclared robustness gate, the conjunction of the four analytic gates, and the additional observed-mapping gate. (B) The internally robust set (16 genes) and the externally supported set (2 genes) do not overlap. (C) Discovery interaction effects against the primary cross-context interaction effects (Heiye response minus Guiwei response, pericarp, 24 h); supported, contradictory, and internally robust genes are highlighted. The diagonal is a visual reference only; equal effects across contexts are not expected.

Figure 4. Frozen pathway and signature evaluation. (A) Discovery normalized enrichment scores for the six frozen pathways; only circadian rhythm passed all internal gates, and none was supported in the primary external evaluation. (B) The frozen signed 206-gene signature across the six non-exploratory external contrasts; filled points denote $q < 0.05$. The quarantined PRJNA1090613 estimate is in Figure S3.


Figure 5. Conditional transcript usage. (A) Discovery gate and event counts. (B) External follow-up of the 225 discovery events, separated by prespecified study role.

Figure 6. Orthogonal evidence and final tiers. (A) Annotation status of the 206 QC-retained genes. (B) Registered promoter-motif robustness and the small-RNA reference gate. (C) Deterministic final tiers over all 268 frozen entities.


1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60




536x472mm (118 x 118 DPI)



549x498mm (118 x 118 DPI)



549x485mm (118 x 118 DPI)



543x472mm (118 x 118 DPI)


1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
50
51
52
53
54
55
56
57
58
59
60

A Conditional DTU discovery




Stage-wise q values and signed usage-interaction differences

B External DTU follow-up remained separated by study role



No primary event was supported; exploratory outcomes remain non-confirmatory.

549x375mm (118 x 118 DPI)



572x427mm (118 x 118 DPI)

Figure S1





Figure S1. Replicate-level normalized counts for the two cross-context-supported genes, twelve Tier B genes, and two legacy highlights. Points are individual deposited libraries; bars show cultivar-treatment medians.

Figure S2



Dashed line: 80% power. Dotted line: registered $|\log_2FC|$ threshold. Parametric power is conditional on fitted means, dispersions, deposited-library independence, and the simulated non-null prevalence.

Figure S2. Parametric detection probability for cultivar-by-infection effects under genome-wide discovery and candidate-family external adjustment. Curves show the overall result and mean-expression quartiles; the dashed line marks 80% power.

Figure S3




Figure S3. Quarantined PRJNA1090613 signed-signature estimate, retained as exploratory rather than confirmatory evidence.